

## Intersection

MovieStar(name,address,gender,birthDate)  
MovieExec(name,address,cert#,netWorth)

Query: Names and addresses of all female movie stars who are also movie executives with netWorth over 10.000.000

First, find female movie stars

```
SELECT name, address  
FROM MovieStar  
WHERE gender = 'F';
```

Next, find movie executives with netWorth over 10.000.000

```
SELECT name, address  
FROM MovieExec  
WHERE netWorth > 10000000;
```

Then take the intersection (note that the schemas are the same)

```
(SELECT name, address  
FROM MovieStar  
WHERE gender = 'F')  
    INTERSECT  
(SELECT name, address  
FROM MovieExec  
WHERE netWorth > 10000000);
```

Difference: Movie Stars who are not Movie Executives

```
(SELECT name, address FROM MovieStar)
  EXCEPT
(SELECT name, address FROM MovieExec);
```

What is wrong with this?

```
(SELECT * FROM MovieStar)
  EXCEPT
(SELECT * FROM MovieExec);
```

## Union

Movies(title,year,length,genre,studioName,producerC#)  
StarsIn(movieTitle, movieYear, starName)

We might have movies in StarsIn that are missing from Movies and vice versa. Find all movies in the database

```
(SELECT title, year FROM Movies)
UNION
(SELECT movieTitle, movieYear FROM StarsIn);
```

Doesn't work, as the schemas are different. Need to rename:

```
(SELECT title, year FROM Movies)
UNION
(SELECT movieTitle AS title, movieYear AS year
FROM StarsIn);
```

## Note on duplication

- SQL uses bag semantics
- Duplicates are always preserved (remember Disney example?)
- Exception: INTERSECT, UNION and EXCEPT use set semantics as default
- If you want to preserve duplicates, you must use INTERSECT ALL, UNION ALL and EXCEPT ALL instead

## Eliminating duplication

Keyword DISTINCT

Movies(title,year,length,genre,studioName,producerC#)

```
SELECT studioName  
FROM Movies;
```

Gives us a list of Studios with many repetitions

```
SELECT DISTINCT studioName  
FROM Movies;
```

Gives us each Studio only once

## Aggregation

The basic concepts are similar to those that we saw for the operator  $\gamma$

Keyword is GROUP BY

Movies(title,year,length,genre,studioName,producerC#)

Find, for each studio, the sum of the lengths of the movies produced by this  
Studio

```
SELECT studioName, SUM(length) AS totalLength  
FROM Movies  
GROUP BY studioName;
```

After using a GROUP BY clause, *only* the attributes in this clause can be used for the selection.

Each StudioName is then associated with a set of all titles, a set of all years, etc, and these can only be accessed using aggregation

Aggregates are AVG, SUM, COUNT, MAX, MIN

The following queries are legal (though don't make much sense. . . ).



```
SELECT studioName, SUM(length) AS totallength,  
       MAX(title) AS maxtitle,  
       AVG(year) AS avgyear, SUM(producerC#) AS sumprod  
FROM Movies  
GROUP BY studioName;
```

```
SELECT studioName, title, SUM(length) AS totallength,  
       AVG(year) AS avgyear  
FROM Movies  
GROUP BY studioName, title;
```

GROUP BY attributes *may* be used in the SELECT-clause, but don't have to be

On the other hand, this is not legal:

```
SELECT studioName, title, SUM(length) AS totallength,  
       AVG(year) AS avgyear  
FROM Movies  
GROUP BY studioName;
```

After using GROUP BY, title can only be used with an aggregate

We can use DISTINCT with aggregation.

Suppose that the movies by Disney have lengths 100, 100, and 120.

Then this query will contain a tuple ('Disney', 320)

```
SELECT studioName, SUM(length) AS totalLength  
FROM Movies  
GROUP BY studioName;
```

while this will contain a tuple ('Disney', 220)

```
SELECT studioName, SUM(DISTINCT length) AS totalLength  
FROM Movies  
GROUP BY studioName;
```

## Another example

Find the total length of film for each producer (by name)

We need to use two relations, one for the film information, and one to find the producer's name

After taking the join, we can proceed with the aggregation on that single (new) relation

```
Movies(title,year,length,genre,studioName,producerC#)  
MovieExec(name,address,cert#,netWorth)
```

```
SELECT name, SUM(length) AS totalLength  
FROM MovieExec, Movies  
WHERE producerC# = cert#  
GROUP BY name;
```

Note use of the WHERE-clause. This is applied first to create the join, then the GROUP BY is applied

## More on selection

Suppose we want to evaluate the previous query only for producers with netWorth more than 10.000.000

```
SELECT name, SUM(length) AS totalLength
FROM MovieExec, Movies
WHERE producerC# = cert# AND netWorth > 10000000
GROUP BY name;
```

What is we want to evaluate the query only for producers who made at least one movie before 1930?

```
SELECT name, SUM(length) AS totalLength
FROM MovieExec, Movies
WHERE producerC# = cert# AND year < 1930
GROUP BY name;
```

Why doesn't this work?

## HAVING clause

This computes the total length of movies made before 1930, while we want to include later movies in the sum

Our condition is actually a condition on each group. We want the group of years corresponding to a producer to contain at least one date before 1930

We could do this with two queries. First, define R as

```
SELECT name, SUM(length) AS totalLength, MIN(year) AS minYear
FROM MovieExec, Movies
WHERE producerC# = cert#
GROUP BY name;
```

And then select those tuples from R with minYear less than 1930

- Requires two queries (we'll see later how to do this)
- Rather complicated.

SQL provides the HAVING clause so that we can do this in one query

The HAVING clause is a condition on aggregates

```
SELECT name, SUM(length) AS totalLength
FROM MovieExec, Movies
WHERE producerC# = cert#
GROUP BY name
HAVING MIN(year) < 1930;
```

## Subqueries

- An SQL query can use subqueries
- Can be used to make the structure of a complex query clearer

Subqueries can produce

- Single values, to be used in WHERE clauses
- Relations, that can be used in WHERE-clauses and FROM-clauses (with a tuple variable)



## Subqueries with a single value

Single value: called a *scalar*

User must be sure that such a subquery will only produce a single value

Recall:

```
SELECT name
FROM Movies, MovieExec
WHERE title = 'Star Wars' AND year=1977
      AND producerC# = cert#;
```

If we knew producerC# (producer code of *Star Wars*) was 1100, we could write

```
SELECT name
FROM MovieExec
WHERE cert# = 1100;
```

We can find producerC# with the query

```
SELECT producerC#  
FROM Movies  
WHERE title = 'Star Wars' and year = 1977;
```

Furthermore, since (title,year) is a key for Movies, we know that this query produces a single value, so we can write

```
SELECT name  
FROM MovieExec  
WHERE cert# = (  
    SELECT producerC#  
    FROM Movies  
    WHERE title = 'Star Wars' and year = 1977);
```

Note that the schema of the subquery doesn't have to match

What happens if we omit year=1977?

## Conditions on relations

What if the subquery does not produce a scalar?

In this example, it will be an error.

Instead of  $=$ , we can use other conditions

For now, assume that  $R$  is a unary relation

- $\text{EXISTS } R$ : There exists a tuple in  $R$ , i.e.,  $R$  is not empty
- $s \text{ IN } R$ : Holds if  $s$  is one of the values in  $R$
- $s \text{ NOT IN } R$ : Holds if  $s$  is not equal to any of the values in  $R$
- $s > \text{ALL } R$ : Holds if  $s$  is greater than all the values in  $R$
- $s > \text{ANY } R$ : Holds if  $s$  is greater than some value in  $R$

We can also use  $>$ ,  $<=$ ,  $>=$ ,  $=$ ,  $<>$ .

What do  $s=\text{ALL } R$ ,  $s=\text{ANY } R$ ,  $s<>\text{ALL } R$ ,  $s<>=\text{ANY } R$  mean?

## Another example

Find the producers of Harrison Ford's movies

Step 1: Find the title and year of these movies

```
SELECT movieTitle, movieYear  
FROM StarsIn  
WHERE starName = 'Harrison Ford';
```

Results include (*Star Wars*, 1977) and (*Raiders of the Lost Ark*, 1981)

Step 2: Find the producer codes for each of these.

```
SELECT producerC#  
FROM Movies  
WHERE title = 'Star Wars' and year = 1977;
```

and

```
SELECT producerC#  
FROM Movies  
WHERE title = 'Raiders of the lost ark' and year = 1981;
```

We get two codes, say 100 and 111.

We could combine these two queries

```
SELECT producerC#  
FROM Movies  
WHERE (title,year) IN  
      (SELECT movieTitle, movieYear  
        FROM StarsIn  
        WHERE starName = 'Harrison Ford');
```

Note use of tuple notation for the IN clause when the relation is binary

Note also that the two schemas don't have to match

Step 3: Find the names of the producers.

```
SELECT name  
FROM MovieExec  
WHERE cert# = 100;
```

Yields the name "Gary Kurtz"

```
SELECT name  
FROM MovieExec  
WHERE cert# = 111;
```

Yields the name "Howard Kazanjian"

We can combine all three parts into

```
SELECT name
FROM MovieExec
WHERE cert# IN
  (SELECT producerC#
   FROM Movies
   WHERE (title,year) IN
     (SELECT movieTitle, movieYear
      FROM StarsIn
      WHERE starName = 'Harrison Ford'));
```

Since the first subquery produces a single value, tuple notation is not required



Without subqueries

```
SELECT name
FROM MovieExec, Movies, StarsIn
WHERE cert# = producerC#
      AND title =  movieTitle
      AND movieYear = year
      AND starName = 'Harrison Ford';
```

## Alternative

```
WITH R AS
  SELECT producerC#
  FROM Movies
  WHERE (title,year) IN
    (SELECT movieTitle, movieYear
     FROM StarsIn
     WHERE starName = 'Harrison Ford')
SELECT name
FROM MovieExec
WHERE cert# IN R;
```

WITH defines  $R$  to be a temporary relation, to be used only in this query

## Correlated subqueries

Much more complicated use of subqueries

`Movies(title,year,length,genre,studioName,producerC#)`

Find all the titles that have been used more than once

Our running example: “King Kong”, with years 1933, 1976, and 2005

If we already knew this title, we could find the years with

```
SELECT year
FROM Movies
WHERE title = 'King Kong';
```

Title has been reused if we get more than one value, here {1933, 1976, 2005}

Alternatively, it has been reused, if we can find some year for this movie (1933 or 1976) which is smaller than some value in this set

In other words, if our relation only contained the *King Kong's*

```
SELECT title
FROM Movies
WHERE year < ANY
    (SELECT year
     FROM Movies
     WHERE title = 'King Kong');
```

In general, introduce a tuple variable *Old* to range over *Movies* and use *Old.title* in the subquery

```
SELECT title
FROM Movies Old
WHERE year < ANY
    (SELECT year
     FROM Movies
     WHERE title = Old.title);
```

Check this gives the same result for *King Kong*. Then check what it does for other movies

## Subqueries in FROM clause

Harrison Ford query again. First create a relation with the producer codes

```
SELECT producerC#  
FROM Movies, StarsIn  
WHERE title = movieTitle AND  
       movieYear = year AND  
       starName = 'Harrison Ford';
```

Then use this in the FROM clause

```
SELECT name  
FROM MovieExec, (SELECT producerC#  
                  FROM Movies, StarsIn  
                  WHERE title = movieTitle AND  
                        movieYear = year AND  
                        starName = 'Harrison Ford'  
                  ) Prod  
WHERE cert# = Prod.producerC#;
```

## Other options

The following keywords can be used as abbreviations for more complex queries, or simply for readability. They correspond to the algebraic operations that we have seen earlier

`Movies CROSS JOIN StarsIn`

`Creates`  $\text{Movies} \times \text{StarsIn}$

Can be used in SQL FROM-clause (with tuple variable, if needed)

## Theta-Join

```
Movies JOIN StarsIn ON  
    title = movieTitle AND year = movieYear;
```

For example,

```
SELECT title, year, length, starName  
FROM Movies JOIN StarsIn ON  
    title = movieTitle AND year = movieYear;
```

Advantage: Separate join conditions from other conditions for readability

```
SELECT title, length, starName  
FROM Movies JOIN StarsIn ON  
    title = movieTitle AND year = movieYear;  
WHERE year = 1977
```

## Natural join

When the attributes for the join are the same

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth

MovieStar NATURAL JOIN MovieExec

creates a relation with attributes

name, address, gender, birthdate, cert#, netWorth



## Outerjoins

Other commands, with the same semantics are the corresponding algebraic operations

```
MovieStar NATURAL FULL OUTER JOIN MovieExec;  
MovieStar NATURAL LEFT OUTER JOIN MovieExec;  
MovieStar NATURAL RIGHT OUTER JOIN MovieExec;
```

If the attribute names differ, we can use

```
MovieStar FULL OUTER JOIN MovieExec ON  
    movieTitle = title AND movieYear = year;  
MovieStar LEFT OUTER JOIN MovieExec ON  
    movieTitle = title AND movieYear = year;  
MovieStar RIGHT OUTER JOIN MovieExec ON  
    movieTitle = title AND movieYear = year;
```

## Database modification

### Database modifications

- Insert: insert new tuples
- Delete: delete tuples from the database
- Modify: modify existing tuples

Why modify? We could always insert and delete tuples.

Some of the reasons include

- Efficiency
- Semantics (example later, when we discuss foreign keys)
- Security: A user may be allowed to modify specific fields only

## Insertions

```
INSERT INTO R(A1,...,An) VALUES (v1,...,vn);
```

First part lists the names of the columns in the order that corresponds to the values (remember that the order of the columns has no meaning)

```
INSERT INTO StarsIn(movieTitle, movieYear, starName)
VALUES ('The Maltese Falcon', 1942,
        'Sydney Greenstreet');
```

But: If we use the order of attributes from the CREATE TABLE command, we can omit the attributes here

```
INSERT INTO StarsIn VALUES
('The Maltese Falcon', 1942, 'Sydney Greenstreet');
```

What if name is a key for a relation and we try to insert a tuple with a value that is already in the relation?

The system should give an error message and refuse to update the relation

Studio(name,address, presC#)

```
INSERT INTO Studio(name) VALUES ('MGM');
```

What happens to the other columns? They become NULL (if allowed, otherwise update disallowed)

We get the tuple ('MGM', NULL, NULL)

Movies(title,year,length,genre,studioName,producerC#)

Studio(name,address, presC#)

Suppose we want to insert in Studio those studio names that are missing from Movies (and only these studios). Other attribute values will have to be NULL

```
INSERT INTO Studio(name)
  SELECT DISTINCT studioName
  FROM Movies
 WHERE studioName NOT IN
    (SELECT name
     FROM Studio);
```