

Deletion

```
DELETE FROM R WHERE <condition>
```

Removes every tuple satisfying the condition from R

```
DELETE FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon' AND
      movieYear = 1942 AND
      starName = 'Sydney Greenstreet';
```

All three conditions may not be necessary to identify a tuple (e.g., just specify key):

```
DELETE FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon' AND
      movieYear = 1942;
```

Deleting several tuples at the same time

```
DELETE FROM MovieExec  
WHERE netWorth < 10000000
```

Removes all tuples for which the value of netWorth is less than 10.000.000

Updates

Update: Modifies part of a tuple

```
UPDATE MovieExec
SET name = 'Pres.' || name
WHERE cert# IN
      (SELECT presC#
       FROM Studio);
```

name on left of SET clause is the new value; name on right is the old one

Tables

Three types

- Stored relations or tables
- Views. Created when needed. To be studied later
- Temporary tables (e.g., WITH clause)

Data Types

- Character strings of fixed or varying length
 - CHAR(n): fixed-length string of up to n characters
 - VARCHAR(n): string of up to n characters
- Bit strings
 - BIT(n): length n
 - BIT VARYING(n): up to length n
- BOOLEAN: Boolean strings. Values can be TRUE, FALSE, UNKNOWN (to be explained later)
- INT or INTEGER: integers (also SHORTINT)

Types, continued

- Floating point

- FLOAT, REAL
- DOUBLE PRECISION
- DECIMAL(*n*,*d*). For example 0123.45 is of type DECIMAL(6,2)

DATE and TIME. Special form of strings:

- DATE '1948-05-14'
- TIME '15:00:02.5'

Table declarations

```
CREATE TABLE Movies (  
    title      CHAR(100),  
    year       INT,  
    length     INT,  
    genre      CHAR(10),  
    studioName CHAR(30),  
    producerC# INT  
);
```

```
CREATE TABLE MovieStar (  
    name       CHAR (30),  
    address    VARCHAR(255),  
    gender     CHAR(1),  
    birthdate  DATE  
);
```

Modifying tables

- DROP TABLE R;

Deletes table *and all of its tuples*

- Modify table. Syntax:

- ALTER TABLE

- Name of relation

- Options. Two are:

- ADD followed by attribute name (and type)
- DROP followed by attribute name

Example

- `ALTER TABLE MovieStar ADD phone CHAR(16);`
- MovieStar now has five attributes
- What happens to the tuples?
- Normally, the value will be set to NULL (but one can specify a different default)
- Another example:
`ALTER TABLE MovieStar DROP birthdate`
deletes the birthdate attribute (the values are lost)

Default values

- Default value (instead of NULL)
- ALTER TABLE MovieStar ADD phone CHAR(16) DEFAULT 'unlisted';
- Other examples

gender CHAR(1) DEFAULT '?'

birthdate DATE DEFAULT DATE '0000-00-00'

Keys

Two ways to declare:

- Part of attribute declaration (single-attribute keys only)
- Separate statement

Two types of declarations:

- PRIMARY KEY
- UNIQUE

Semantics: No two different tuples can have the same values for the key attribute(s)

PRIMARY KEY vs. UNIQUE

- Every relation must have exactly one PRIMARY KEY
- A relation may have one or more UNIQUE statements
- PRIMARY KEY attributes cannot be NULL
- UNIQUE attributes may be. Two different tuples may have NULL values for UNIQUE attributes
- An implementation usually creates an efficient access method for PRIMARY KEY attributes

Examples

```
CREATE TABLE MovieStar (  
    name      CHAR(30) PRIMARY KEY,  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE  
);
```

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE,  
    PRIMARY KEY (name)  
);
```

Two-attribute keys

```
CREATE TABLE Movies (  
    title      CHAR(100),  
    year       INT,  
    length     INT,  
    genre      CHAR(10),  
    studioName CHAR(30),  
    producerC# INT,  
    PRIMARY KEY (title, year)  
);
```

Constraints

- Constraints are restrictions on the contents of the database
- The system guarantees their validity
- Attempts to violate them will be rejected by the system
- Two types of constraints: Constraints on a single relation (checked only when modifying this relation) and constraints on the whole database (checked on any modification)
- Examples we have seen: PRIMARY KEY and UNIQUE

Example of PRIMARY KEY and UNIQUE constraints

```
CREATE TABLE Student (  
    name          CHAR(30),  
    codFiscale    CHAR(16),  
    matricola     INT,  
    PRIMARY KEY (codFiscale),  
    UNIQUE (matricola)  
);
```

- Also allowed: “matricola INT UNIQUE” provided UNIQUE refers to a single attribute
- PRIMARY KEY and UNIQUE are quite similar
- One (and only one) PRIMARY KEY is required for every relation
- Any number of UNIQUE constraints are allowed
- PRIMARY KEY is usually used by the system to construct an efficient access mechanism
- PRIMARY KEY attribute(s) can never be NULL; UNIQUE attributes can be

Foreign Keys: Second most important class of constraint

Studio(name, address, presC#)

MovieExec(name, address, cert#, netWorth)

presC# can be seen as a reference to a tuple in MovieExec if

- cert# is declared as a key or as UNIQUE
- For every value of presC# in Studio, there is a tuple in MovieExec with this value for cert#
- Exception: when presC# is NULL

Two formats

```
CREATE TABLE Studio(  
    name      CHAR(30) PRIMARY KEY,  
    address   VARCHAR(255),  
    presC#    INT REFERENCES MovieExec(cert#)  
)
```

```
CREATE TABLE Studio(  
    name      CHAR(30) PRIMARY KEY,  
    address   VARCHAR(255),  
    presC#    INT,  
    FOREIGN KEY (presC#) REFERENCES MovieExec(cert#)  
)
```

Foreign keys can also contain more attributes; in this case only the second syntax is allowed

FOREIGN KEY syntax does not require that cert# be a key, but it is always used in this way

Modifications

What happens when we try to modify one of the two relations?

Modifying Studio

- Deletion: Nothing can go wrong
- Insertion: If we insert a tuple with non-NULL presC# field which is not present in MovieExec, we must reject the insertion
- Same with modification of the presC# of an existing tuple. All other updates (i.e., modifications of other attributes) are OK

Modifying MovieExec

- Inserting tuples: No problem
- Deleting a tuple whose cert# exists in Studio: A problem
- Same problem with modifying a cert# field that exists in Studio

Policies

In the case of changes to MovieExec, we have several choices

- Reject the modification
- Cascade
 - In the case of deletion, delete all the tuples in Studio that refer to this tuple. There might be tuples in another relation that refer to these, in that case delete them, and so on
 - In the case of update, change the values of cert# in Studio that refers to this tuple to the new value
- Change the cert# fields that refer to the modified tuple to NULL

SQL syntax

```
CREATE TABLE Studio(  
    name      CHAR(30) PRIMARY KEY,  
    address   VARCHAR(255),  
    presC#    INT REFERENCES MovieExec(cert#)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
)
```

Constraints on a single relation

Single relation constraints can be on a single attribute or a single tuple

One example: NOT NULL says that the attribute is not allowed to be equal to NULL. For example

```
presC#  INT REFERENCES MovieExec(cert#) NOT NULL
```

Note that in this case ON DELETE SET NULL cannot be used

CHECK constraints

Two types:

- conditions on individual attributes
- conditions on individual tuples

Both are part of table declaration. First type *can* be part of attribute declaration

- In Studio

```
presC#  INT REFERENCES MovieExec(cert#)
        CHECK (presC# >= 1000000),
```

- In MovieStar

```
gender CHAR(1) CHECK (gender IN ('F','M')),
```

Syntax

- The condition in CHECK is a SQL condition with BOOLEAN result
- Condition can use other relations, but this is usually a bad idea
- Attempt to define foreign key using CHECK

```
presC# INT CHECK  
    (presC# IN (SELECT cert# FROM MovieExec)),
```

- What is wrong with this?

When constraints are checked

- CHECK are tested when the relation in which they are declared is changed
- Therefore, this condition is not tested when MovieExec is changed
- Foreign keys are tested when either relation changes

Tuple-based constraints: constraints on whole tuple

```
CHECK (gender = 'F' OR name NOT LIKE 'Ms. %')
```

Syntax for single-attribute conditions

Two possibilities

```
CREATE TABLE MovieStar(  
    name          CHAR(30)  PRIMARY KEY,  
    address       VARCHAR (255),  
    gender        CHAR (1)  CHECK (gender in ('F', 'M')),  
    birthdate     DATE  
)
```

or

```
CREATE TABLE MovieStar(  
    name          CHAR(30)  PRIMARY KEY,  
    address       VARCHAR (255),  
    gender        CHAR (1),  
    birthdate     DATE,  
    CHECK (gender in ('F', 'M'))  
)
```

First is better, as it only needs to be checked when that attribute is changed

Modifying constraints

- In Studio

```
presC# INT CHECK (presC# >= 1000000),
```

- We want to change the constraint. This can only be done if the constraint is declared with a name

```
presC# INT CONSTRAINT SixDigits  
CHECK (presC# >= 1000000)
```

- We can then write

```
ALTER TABLE Studio DROP CONSTRAINT SixDigits
```

```
ALTER TABLE Studio ADD CONSTRAINT SevenDigits  
CHECK (presC# >= 10000000)
```

Modification of constraints

- To modify a constraint it must be given a name when declared

- Examples

```
name CHAR(30) CONSTRAINT NameIsKey PRIMARY KEY
```

```
gender CHAR(1) CONSTRAINT NoAndro CHECK (gender IN (F>, >M)),
```

- Delete constraint

```
ALTER TABLE MovieStar DROP CONSTRAINT NameIsKey;
```

- Add constraint

```
ALTER TABLE MovieStar ADD CONSTRAINT NameIsKey PRIMARY KEY (name);
```

Assertions

- Outside the declarations of tables
- Boolean queries, that must always be true
- Should be checked after every update
- Example

No studio can have a president with netWorth less than 10.000.000
(but other executives could)

```
CREATE ASSERTION RichPres CHECK
  (NOT EXISTS
    (SELECT Studio.name
     FROM Studio, MovieExec
     WHERE presC# = cert# AND netWorth < 10000000));
```

Another example

Aggregation

Student(matricola, name, course)

No course can have more than 50 students

```
CREATE ASSERTION numeroChiusa CHECK
  (NOT EXISTS
    (SELECT course
     FROM student
     GROUP BY course
     HAVING COUNT(matricola) > 50))
```

Triggers

- Triggers are commands that are executed when something “happens” to the database
- We will look at triggers on updates
- Trigger has a condition that is checked when the database is changed
- If the condition holds, the trigger is activated

Example

- Trigger that does not allow reducing the value of netWorth
- Why can't we use a constraint?
- Intuition: After a “bad” update, restore the old value

```
CREATE TRIGGER NetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN (OldTuple.netWorth > NewTuple.netWorth)
    UPDATE MovieExec
    SET netWorth = OldTuple.netWorth
    WHERE cert# = NewTuple.cert#;
```



```
CREATE TRIGGER NetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN (OldTuple.netWorth > NewTuple.netWorth)
    UPDATE MovieExec
    SET netWorth = OldTuple.netWorth
    WHERE cert# = NewTuple.cert#;
```

- Declare trigger and give it a name
- When to apply trigger. Options:
 - AFTER, BEFORE, INSTEAD OF
 - INSERT, DELETE, UPDATE
 - OF: Only makes sense (and is optional) for UPDATE
 - ON: relation name

Accessing data

```
CREATE TRIGGER NetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN (OldTuple.netWorth > NewTuple.netWorth)
    UPDATE MovieExec
    SET netWorth = OldTuple.netWorth
    WHERE cert# = NewTuple.cert#;
```

- REFERENCING followed by variable names for OLD (before modification) and NEW (after)
- OLD only allowed for update and delete; NEW for update and insert
- Options: ROW (tuple) and TABLE (entire relation)

```
CREATE TRIGGER NetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN (OldTuple.netWorth > NewTuple.netWorth)
    UPDATE MovieExec
    SET netWorth = OldTuple.netWorth
    WHERE cert# = NewTuple.cert#;
```

- FOR EACH ROW or FOR EACH STATEMENT (update statement)
- WHEN optional
- SQL statement, or list of SQL statements (with BEGIN and END)

Another example

- Constraint on Movie Exec: Don't allow average of netWorth to be less than 500.000
- Cannot do this by checking each tuple
- Needs to be done for the complete update statement
- This example could also be done by an assertion. With a more complicated trigger, we could also say how to fix the update
- We describe a trigger for updates. Similar triggers must also be written for insertion and deletion

```
CREATE TRIGGER AvgNetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD TABLE AS OldStuff,
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (500000 > (SELECT AVG(netWorth) FROM MovieExec))
BEGIN
    DELETE FROM MovieExec
    WHERE (name, address, cert#, netWorth) IN NewStuff;
    INSERT INTO MovieExec
        (SELECT * FROM OldStuff);
END;
```

BEFORE triggers

- Insert tuples into Movies(title,year,length,genre,studioName,producerC#)
- What if we don't know the year?
- Cannot insert a NULL value, since year is part of the key
- Check for NULL and change the value if needed
- We use 1915, which could be done with default values, but more complicated rules can be implemented with triggers

```
CREATE TRIGGER FixYearTrigger
BEFORE INSERT ON Movies
REFERENCING
    NEW ROW AS NewRow
    NEW TABLE AS NewStuff
FOR EACH ROW
    WHEN NewRow.year IS NULL
    UPDATE NewStuff SET year = 1915;
```

Views

- View: Defined by a query
- Creates a virtual table: Can be used by queries *as though* it was another relation
- But the relation is not physically constructed

Used for

- Convenience
- Security: Allow users only to access views, which hide the rest of the data

Why not create physical relations?

- Save space
- No need to update them

Example

```
CREATE VIEW ParamountMovies AS
  SELECT title, year
  FROM Movies
  WHERE studioName = 'Paramount';
```

User may be permitted access only to this relation, hiding many tuples as well as some attributes

Another example:

```
CREATE VIEW MovieProd
  SELECT title, name
  FROM Movies, MovieExec
  WHERE producerC# = cert#;
```

Saves the user from writing a join to find the producer of a movie

Use of Views

Views can be used *as though* they were normal relations

Find movies made by Paramount in 1979

```
SELECT title
FROM ParamountMovies
WHERE year = 1979;
```

One can use relations and views in the same query.

Find all the names of Stars in Paramount movies

```
SELECT starName
FROM ParamountMovies, StarsIn
WHERE title = movieTitle AND year = movieYear;
```

This is equivalent to

```
SELECT starName
FROM (SELECT title, year
      FROM Movies
      WHERE StudioName = 'Paramount'
      ) Pm, StarsIn
WHERE Pm.title = movieTitle AND Pm.year = movieYear;
```

Another example: Rename attributes in View

```
CREATE VIEW MovieProd(moveTitle, prodName) AS
  SELECT title, name
  FROM Movies, MovieExec
  WHERE producerC# = cert#;
```

VIEW and WITH

- Views are always available, until they are deleted
- WITH creates a relation than is only available during execution of a query