

Embedded SQL

- Use SQL from inside a programming language
- Any language: C++, Java, etc
- We look at general, non-language specific aspects
- Why use embedded SQL?
 - Formatting (e.g. HTML)
 - Express queries not expressible in SQL (complex recursion)
 - Database integration, extensions etc.

Basic idea

- SQL statements used as subprocedures
- Main problem: Arguments and results in C++ (for example) have types
- Basic type or relation doesn't exist in C++, and SQL doesn't have pointers etc.

Basic concepts

- EXEC SQL
- Followed by SQL statement or statements (surrounded by BEGIN and END)
- Shared variables. Prefix ':' in SQL, without ':' in C++
- Special variable SQLSTATE (array of 5 characters)

Variable declarations

```
EXEC SQL BEGIN DECLARE SECTION;  
    char studioName[50], studioAddr[256];  
    char SQLSTATE[6];  
EXEC SQL END DECLARE SECTION;
```

Declares variables that can be used both in C++ program and in SQL calls

We describe a function that asks user for values of attributes and inserts them into Studio (we omit non-SQL parts of the code)

Insertion example

```
void getStudio() {  
  
EXEC SQL BEGIN DECLARE SECTION;  
    char studioName[50], studioAddr[256];  
    char SQLSTATE[6];  
EXEC SQL END DECLARE SECTION;  
  
/* print request that studio name and address be entered and read  
response into variables studioName and studioAddr */  
  
EXEC SQL INSERT INTO Studio(name, address)  
    VALUES (:studioName, :studioAddr);  
}
```

This function has no output, so integration with SQL is easy

Queries with output

- Single-tuple SELECT: Direct communication with shared variables
- General queries: *cursors*
- Example: Read studio name, and print president's netWorth

Example

```
void printNetWorth {  
  
    EXEC SQL BEGIN DECLARE SECTION;  
        charstudioName [50]; int presNetWorth;  
        char SQLSTATE[6];  
    EXEC SQL END DECLARE SECTION;  
  
    /* print request that studio name be entered, read response into  
       studioName */  
  
    EXEC SQL SELECT netWorth  
        INTO :presNetWorth  
        FROM Studio, MovieExec  
        WHERE presC# = cert# AND  
            Studio.name = :studioName;  
  
    /* check that SQLSTATE has all 0s (meaning: no error) and if so,  
       print the value of presNetWorth */  
  
}
```

Cursors

- Cursor: A pointer that iterates over tuple in a relation
- Relation can be base relation or result of query
- Declare cursor

```
EXEC SQL DECLARE <cursor name> CURSOR FOR <query>
```

<query> is relation name or SQL query

- Initialize cursor

```
EXEC SQL OPEN <cursor name>
```

Cursor is now ready to get the first tuple

- Opening is separate from declaration so that we can examine the relation more than once

Fetching tuples

- Gets the next tuple of the relation over which the cursor ranges
- Syntax:

```
EXEC SQL FETCH FROM<cursor name> INTO <list of variables>
```

- One variable in the list for each attribute of the relation
- Reads the tuple and advances to the next tuple (even if it doesn't exist)
- If we are at the last tuple, return nothing and set SQLSTATE to 02000 (no need to memorize these codes...)

```

void worthRanges() {

    int i, digits, counts[15];
    EXEC SQL BEGIN DECLARE SECTION;
        int worth;
        char SQLSTATE[6];
    EXEC SQL END DECLARE SECTION;
    EXEC SQL DECLARE execCursor CURSOR FOR
        SELECT netWorth FROM MovieExec;

    EXEC SQL OPEN execCursor;
    for (i=1; i<15; i++) counts[i] = 0;
    while(1) {
        EXEC SQL FETCH FROM execCursor INTO :worth;
        if (NO_MORE_TUPLES) break;
        digits = 1;
        while((worth /= 10) > 0) digits++;
        if(digits <= 14) counts[digits]++;
    }
    EXEC SQL CLOSE execCursor;
    for(i=0; i<15; i++)
        printf ("digits = %/d: number of execs = %d\n",
            i, counts[i]);
}

```

Modifcations

- We can modify the tuple at the current location of the cursor
- Syntax is standard SQL with the condition `WHERE CURRENT`
- Example: examine value of `netWorth`:
 - If it is less than 1.000, delete the tuple
 - Otherwise, double the value

```

void changeWorth() {

    EXEC SQL BEGIN DECLARE SECTION;
        int certNo, worth;
        char execName [31], execAddr[256], SQLSTATE[6];
    EXEC SQL END DECLARE SECTION;
    EXEC SQL DECLARE execCursor CURSOR FOR MovieExec;

    EXEC SQL OPEN execCursor;
    while(1) {
        EXEC SQL FETCH FROM execCursor INTO :execName,
            :execAddr, :certNo, :worth;
        if (NO_MORE_TUPLES) break;
        if (worth < 1000)
            EXEC SQL DELETE FROM MovieExec
                WHERE CURRENT OF execCursor;
        else
            EXEC SQL UPDATE MovieExec
                SET netWorth = 2 * netWorth
                WHERE CURRENT OF execCursor;
    }
    EXEC SQL CLOSE execCursor;
}

```

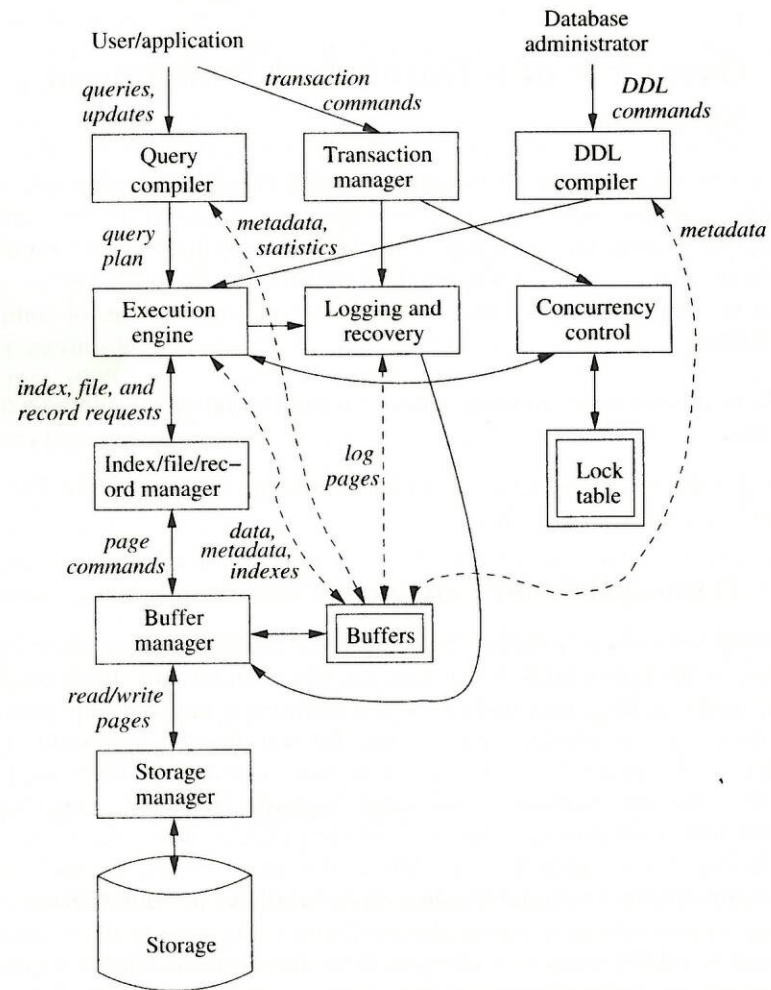
Dynamic SQL

- So far, the SQL statement is part of the C++ program
- What if we want to change the SQL query?
- Basic idea: Use a string variable to contain the SQL statement
- 'SQL variable': variable that holds an interpreted SQL string
- Then use special primitives to convert the string to SQL and execute
- Two statements
 - EXEC SQL PREPARE V FROM <expression> Analyse the SQL statement and find a good strategy to evaluate
 - EXEC SQL EXECUTE V: Execute the query
- We can prepare a query, and then evaluate it multiple times
- Can be combined in one statement EXEC SQL EXECUTE IMMEDIATE <expression>

Example: Read query from user input and execute it

```
void readQuery() {  
  
    EXEC SQL BEGIN DECLARE SECTION;  
        char *query;  
    EXEC SQL END DECLARE SECTION;  
  
    /* prompt user for a query, allocate space (e.g.,  
        use malloc) and make shared variable :query point  
        to the first character of the query */  
    EXEC SQL PREPARE SQLquery FROM :query;  
    EXEC SQL EXECUTE SQLquery;  
}
```

SQL implementation



Query evaluation

- Parsing. Convert query into a parse tree
- Convert parse tree into an algebraic expression
- Algebraic optimization
- Physical plan generation

Optimization:

- Best algebraic expression for query evaluation
- Chose among different algorithms for each algebraic operator
- Decide how to pass data between the operations

Data storage

- Data in databases is usually stored on disks, and the cost of disk access is usually the dominant cost
- The most important thing to optimize is the number of disk accesses: only read the part of the database that we need, and try to avoid reading the data more than once, when possible
- We assume a fixed size of disk blocks, and assume that we can only read data in blocks (rather than tuples)

Indexes

- Index: A permanent data structure that lets us find efficiently all the tuples with a particular value, or range of values, for a given attribute or attributes (such as a key)
- The standard data structure, called a *B-tree* is a search tree as you have seen in your data structure course (depth $\log n$)
- There are two ways in which the B-tree differs from those that you have seen:
 - The leaves contain not the data itself but pointers to the blocks containing the data
 - It must be possible to *update* the tree efficiently, and with the B-tree this can be done in average time $\log n$ (technically amortized time, which is somewhat better)

Clustered index

Two types of index

- Clustered: All tuples with a given value for the index attribute are found on the same page (or adjacent one)
- Unclustered
- (Note: The book also discusses *clustering indexes* which are something different)

Important points

- If there are a small number of tuples (e.g. 10) with a given value for the index, we have to retrieve only 1 (or 2) blocks with a clustered index
- We may have to retrieve 10 blocks with an unclustered one
- We can have at most one clustered index per relation
- There is no limit on the number of unclustered indexes, but the cost in space and updating costs must be taken into account
- Range queries (attribute between k_1 and k_2) can also be answered with B-trees (with similar tradeoffs between clustered and unclustered indexes)


Parsing

- Parser analyses the query and converts it into a tree
- Checks for syntax errors
- Resolves ambiguity. Determines which relation each attribute belongs to
- Checks types: arithmetic only on numeric attributes etc.
- Converts view references to parse trees of their definitions
- We do not study this further. The book contains examples of a *grammar* for SQL. Constructing a parser is a standard application of techniques studied in formal languages

Conversion to relational algebra

- We start with two simple examples, one using a view
- We then discuss some aspects of nested queries
- First example: relations $R(A,B)$, $S(B,C)$


```
SELECT A
FROM R, S
WHERE R.B=S.B AND C=4
```



- FROM: Cartesian product
- WHERE: Selection
- SELECT: Projection

Another example:

SELECT DISTINCT A
FROM R, S
WHERE R.B=S.B AND C=4



Query with view

View

```
CREATE VIEW ParamountMovies AS
  SELECT title, year
  FROM Movie
  WHERE studioName = 'Paramount';
```

Query

```
SELECT title
FROM Paramount Movies
WHERE year = '1979';
```

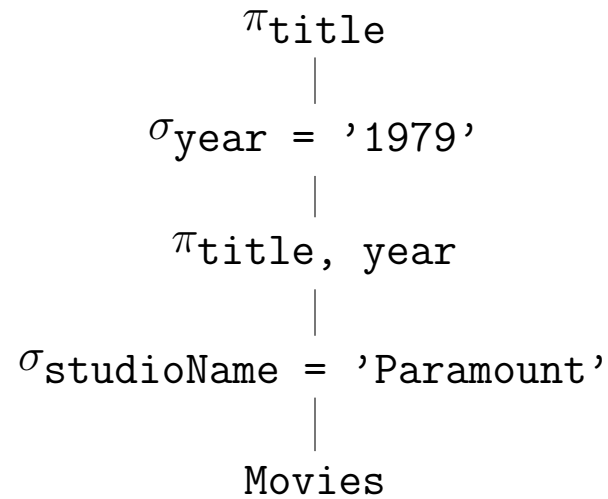
Algebraic expression for view

$$\pi_{\text{title, year}} \mid \sigma_{\text{studioName} = \text{'Paramount'}} \mid \text{Movies}$$

Algebraic expression for query

$$\begin{array}{c} \pi_{\text{title}} \\ | \\ \sigma_{\text{year} = '1979'} \\ | \\ \text{ParamountMovies} \end{array}$$

Combined



Note that the second projection isn't really needed, and should be removed later

Nested queries

- Correlated query: Subquery can refer to attributes in the main query
- Hardest part of the algorithm
- We assume queries are not correlated, so we can convert the subqueries separately from the main query

Queries with EXISTS

```
SELECT A  
FROM R  
WHERE EXISTS Q
```

- Convert Q to algebraic expression \hat{Q}
- $R \times \hat{Q}$ is empty when EXISTS Q is false
- If EXISTS Q is true, $R \times \hat{Q}$ contains all the tuples in R used in the query (care needed with duplicates)
- We get $\pi_A(R \times Q)$
- NOT EXISTS is similar but a bit more complicated

Other subqueries

- Try to convert to queries with EXISTS or NOT EXISTS
- We don't give a formal description, but illustrate with several examples

```
SELECT movieTitle FROM StarsIn
WHERE starName IN (SELECT name
                   FROM MovieStar
                   WHERE birthdate = 1960)
```

Becomes

```
SELECT movieTitle FROM StarsIn
WHERE EXISTS (SELECT name
              FROM MovieStar
              WHERE birthdate = 1960 AND name = starName)
```

Another example

```
SELECT name FROM MovieExec
WHERE netWorth >= ALL (SELECT E.netWorth
                       FROM MovieExec E)
```

Becomes

```
SELECT name FROM MovieExec
WHERE NOT EXISTS (SELECT E.netWorth
                  FROM MovieExec E
                  WHERE netWorth < E.netWorth)
```


Third example

```
SELECT C FROM S
WHERE C IN (SELECT SUM(B) FROM R
GROUP BY A)
```

Becomes

```
SELECT C FROM S
WHERE EXISTS (SELECT SUM(B) FROM R
              GROUP BY A
              HAVING SUM(B) = C)
```

Other primitives

- GROUP BY, HAVING: use γ
- ORDER BY: Use τ
- UNION, INTERSECT, EXCEPT: Use \cup , \cap , $-$

SELECT conditions

- AND: combine conditions
- OR: Use union (careful with duplicates)
- NOT: Use difference

Algebraic optimization

“Pushing” selections

- Selection usually reduces significantly the size of the relations
- Extreme example (but also typical): Selecting for a value of Codice Fiscale reduces a huge relation to a single tuple
- For this reason, we should do selection as soon as possible
- We use algebraic rules that we studied earlier
- For example:


$$\sigma_C(\delta(R)) = \delta(\sigma_C(R))$$

or

$$\sigma_C(R \times S) = \sigma_C(R) \times S$$

provided that C only uses attributes of R


Example




Is first converted to

$$\begin{array}{c} \delta \\ | \\ \pi_A \\ | \\ \sigma_{R.B=S.B} \\ | \\ \sigma_{C=4} \\ | \\ \times \\ / \quad \backslash \\ R \quad S \end{array}$$

We can then push the second selection down the right (not the left) branch to get the query plan



- Query interpreters have very efficient algorithms for join
- Therefore, replace cartesian product with selection by join when possible




- Projections can sometimes be pushed down as well
- The saving is smaller and we do not study this aspect further
- In our example, the projection cannot be pushed down any further (why?)

Another example

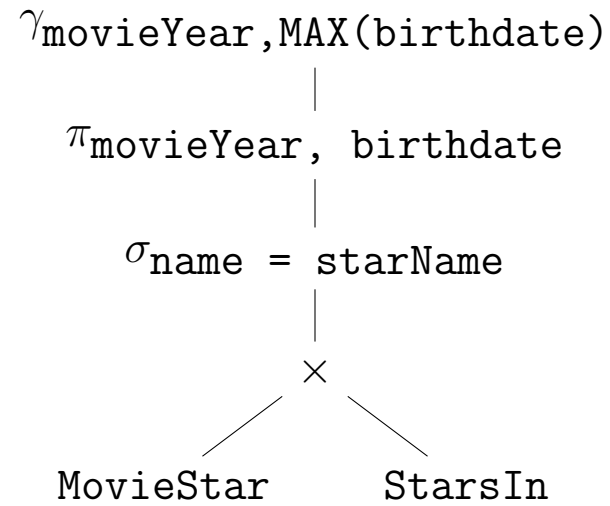
```
SELECT movieYear, MAX(birthdate)
FROM MovieStar, StarsIn
WHERE name = starName
GROUP BY movieYear;
```

Becomes

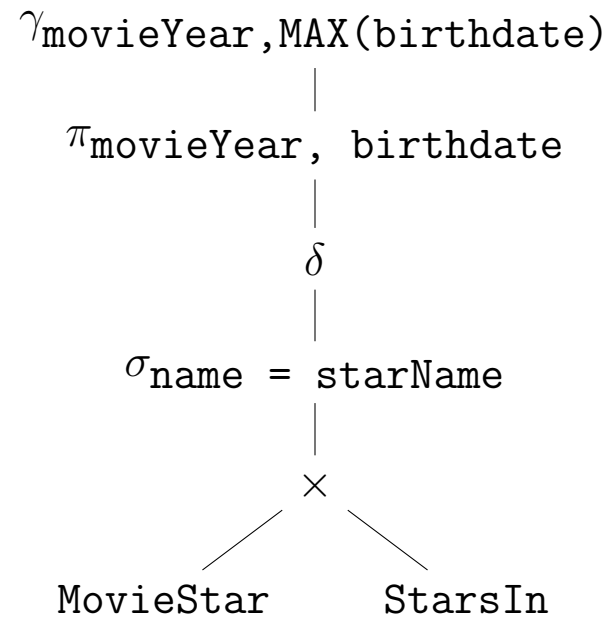


Possible transformations

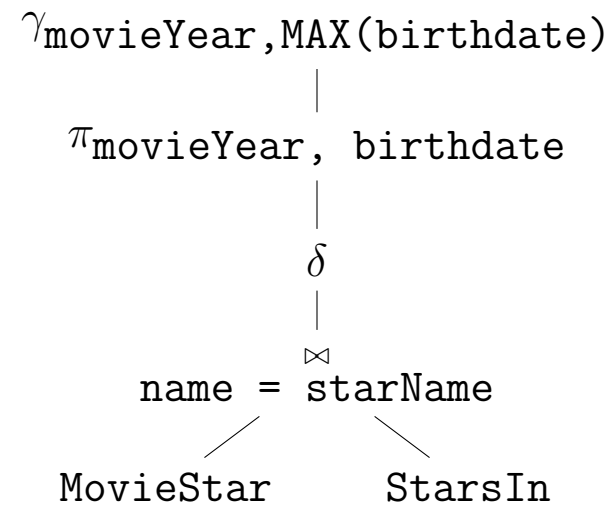
Introduce projection before γ , to remove unneeded attributes



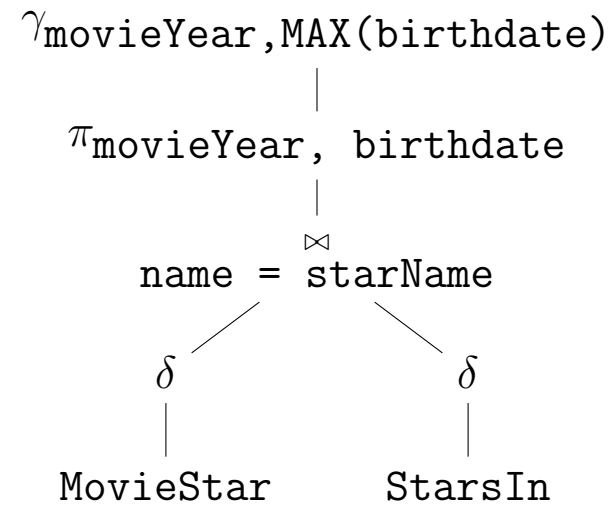
Introduce δ as γ (with MAX) is insensitive to duplicate elimination



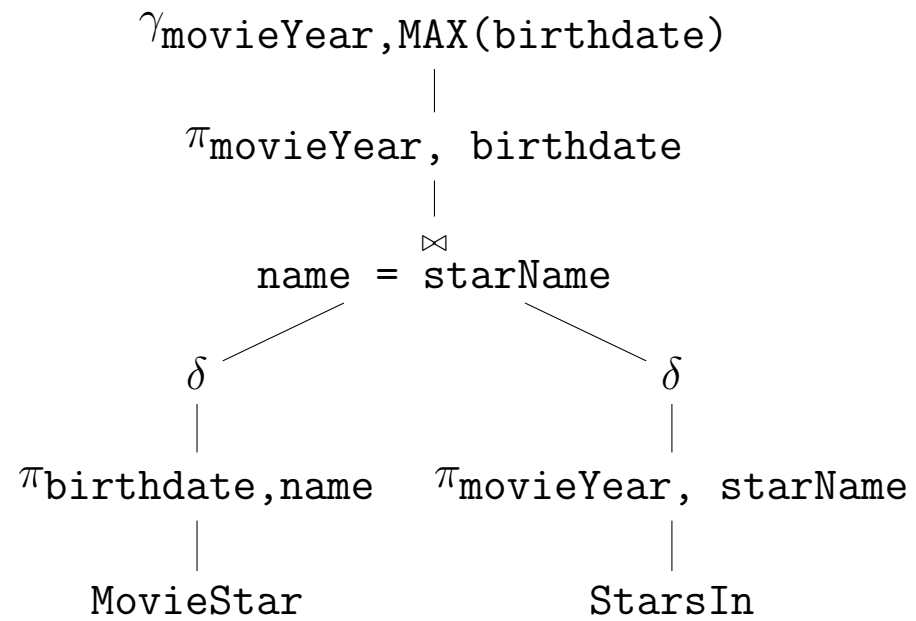
Combine selection and cartesian product into join



Push down δ



Add projections



Which of these algebraic plans are better?

- In general, there is no obvious answer
- The relative efficiency could depend on the data
- We shall study ways of using statistical properties of the data to determine the better plan
- These techniques are not perfect, and will sometimes give a wrong answer
- But errors affect the efficiency only, and not the correctness of the result

A few rules are always better

- Do selections as soon as possible
- Introduce joins instead of cartesian products
- The rule “do selections as soon as possible” is not as simple as it sounds
- The rule applies to the final plan: There are cases where one first moves selections *up* the tree, and then moves them down two branches

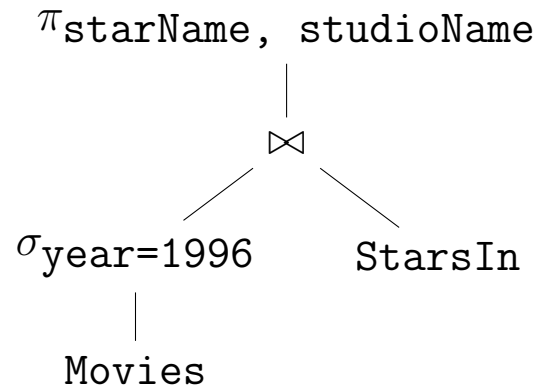
Example:

```
CREATE VIEW MoviesOf1996 AS SELECT *  
  FROM Movies  
  WHERE year = 1996;
```


```
SELECT starName, studioName  
FROM MoviesOf1996 NATURAL JOIN StarsIn;
```

Note that NATURAL JOIN can be directly converted into \bowtie


Algebraic plan for query



Move the selection first up



and then down



Which is the better plan

How do we choose between different evaluation plans?

- Certain rules, such as pushing selections down the tree, are always better
- In other cases, there is no good answer. The choice depends on the data
- We therefore need to know something about statistical properties of the data. This can be gathered by
 - Sampling of the data
 - Keeping track of previous queries
 - Domain-specific rules for the data
 - General rules that do the best possible in the absence of any of the above (this is the approach that we shall take)

Joins


- Join is commutative and associative
- The order in which we combine relations doesn't matter
 - $R \bowtie S = S \bowtie R$
 - $R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$
- But the order can make a huge difference in practice

Example


- $R(A, B)$, $S(B, C)$ and $T(C, D)$
- Suppose that each has 1.000.000 tuples
- Suppose, furthermore, that attribute B is always equal to 1 (this is an extreme case; in practice it is unlikely to be this bad) and that only one value of C in S is equal to only one value of C in T
- If we evaluate $R \bowtie S$ and then join the result to T we create a relation with 10^{12} tuples, and then join some of them to T getting 10^6 tuples
- If, on the other hand, we start with $S \bowtie T$, we get exactly one tuple, and joining it to R is much less expensive

This is an extreme example, but selection of the right join order is still very important in practice


- For now, we postpone this issue
- We represent the join as an operator $\bowtie (R, S, T)$, the join of three (or more) relations
- In theory, we should also give the join conditions, but we only consider natural join
- The tree




becomes



The same is done with other associative and commutative operators such as union and intersection



becomes



Search

- We have to pick the “best” tree, as discussed next
- First, a few words on *which* trees to examine
- For complex queries, examining all trees is unpractical
- There are standard heuristical techniques (such as in AI) for chosing which trees to search
- The techniques for SQL are based on these, and we only discuss search techniques for join order