

## Natural join

- We always use  $R(A, B) \bowtie S(B, C)$
- One-pass algorithms: If  $B(R) \leq M$  the same techniques can be used as for other operations
- Two-pass algorithms: More problematical
- To see why, let's try to do a sort-based 2-pass algorithm for join. The first pass (sorting on  $B$ ) is done as before
- For the second pass, divide (in some way — we'll discuss this later) main memory into two parts, for  $R$  and  $S$ , and assume that the tuples in the first block of each list (for  $R$  and  $S$ ) *all* have the same value  $b$  for  $B$ . Then we cannot perform the join without reading some of  $R$  or  $S$  twice.

Before returning to this example, we'll look at an intermediate class of *nested-loop* or “one-and-a-half” pass algorithms.

In each algorithm, one of the arguments is read once, while the other is read more often

No restrictions on relation size

Tuple-based nested-loop join

```
FOR each tuple s in S DO
  FOR each tuple r in R DO
    IF r and s join to make a tuple t THEN
      output t ;
```

This only reads  $S$  once, but could take  $T(R)T(S)$  iterations!

We'll see later a block-based variant of this algorithm that is much better

But first, an iterator version

```

Open() {
    R.Open();
    S.Open();
    s := S.GetNext();
}

GetNext() {
    REPEAT {
        r := R.GetNext();
        IF (r = NotFound) { /* R is exhausted for
                           the current s */

            R.Close()
            s := S.GetNext();
            IF (s = NotFound) RETURN NotFound
                /* both R and S are exhausted */
            R.Open();
            r := R.GetNext();
        }
    }
    UNTIL (r and s join);
    RETURN the join of r and s;
}

Close() {
    R.Close();
    S.Close();
}

```

## Block-based nest-loop join

Modify the previous algorithm

- Organize access by blocks
- Use as much main memory as we can for  $S$

We assume here that  $S$  is the smaller relation, but will not fit in main memory

$$M < B(S) \leq B(R)$$

```

FOR each chunk of M blocks of S DO BEGIN
  read these blocks into main-memory;
  FOR each block b of R DO BEGIN
    read b into main memory;
    FOR each tuple t of b DO BEGIN
      find the tuples of S in main memory that join with t;
      output the join of t with each of these tuples;
    END;
  END;
END;

```

Note that the inner loop does not involve any I/O

Example:  $B(R) = 1.000$ ,  $B(S) = 500$ ,  $M = 100$

- Outer loop executed 5 times, each time 100 I/O
- Each iteration: Read  $R$  once, 1.000 blocks
- Total time  $5 \times 1.100 = 5.500$

If we switched  $S$  and  $R$ , we iterate 10 times, 600 I/O each time, cost 6.000,  
so starting with the smaller relation makes sense

## Analysis

- Assume  $S$  is the smaller relation
- Outer loop executed  $\frac{B(S)}{M}$  times
- Each iteration reads  $M$  blocks of  $S$  and  $B(R)$  blocks of  $R$
- Total cost  $B(S) + \frac{B(R)B(S)}{M}$
- If  $B(S)$  is much larger than  $M$ , we can approximate this by  $\frac{B(R)B(S)}{M}$
- For reasonably small relations, not much larger than a one-pass join
- For larger relations, much more expensive, but we shall see later that we can use it as a subroutine for small parts of a larger relation

## Sort-based join

- First phase: Sort  $R$  and  $S$  on the join attribute  $B$ .
- Cost:  $4(B(R) + B(S))$ , as we are fully sorting both relations
- Using (for now) only two blocks of main memory, read the first blocks of  $R$  and  $S$
- Find the first value of  $B$  in these blocks. if this value occurs only in one of the relations, delete all these tuples, and continue to the next value
- Read all tuples from both  $R$  and  $S$  that have this value for attribute  $B$ . Assume that they fit in the  $M - 1$  remaining memory blocks
- Output all combinations of these tuples, and delete these tuples in main memory.
- We are left with only two partially filled blocks. Continue from this point

## Example

- $B(R) = 1.000$ ,  $B(S) = 500$ ,  $M = 100$
- Sorting:  $4(1.000 + 500) = 6.000$
- Provided that we never need more than 101 blocks for tuples with the same join value, the last step needs  $1.000 + 500 = 1.500$  I/O
- Total: 7.500 I/O
- Nested-loop join took 5.500 which looks better
- But this is just because both relations are quite small. With  $B(R) = 10.000$  and  $B(S) = 5.000$ , the new algorithm is much better



## Hash-join

Similar to other hashing based algorithms

- Hash based on join attributes
- All tuples that join must be in the same bucket for  $R$  and the same bucket for  $S$
- As long as these two buckets take less than  $M$  blocks (together) we can read them into main memory and join
- Complexity is similar to the previous algorithm

## Index-based joins

- Idea: Use index on join attribute to find tuples that join
- Performance is usually not very good

Assume  $S$  has an index on  $B$

- Examine each block of  $R$
- Consider a tuple  $t$  in this block
- Let  $t_B$  be the value of the  $B$ -attribute
- Use the index to find all the tuples in  $S$  with this value for  $B$
- Output the join of these tuples with  $r$

## Analysis

- Accessing  $R$ :  $B(R)$  I/O
- For each tuple of  $R$ , on average  $\frac{T(S)}{V(S,B)}$  have the same value for attribute  $B$
- This must be done for each tuple of  $R$ , yielding cost  $\frac{T(R)T(S)}{V(S,B)}$
- Total  $B(R) + \frac{T(R)T(S)}{V(S,B)}$

## Example

The use of  $T(R)$  and  $T(S)$  already should show that this is very expensive

- $B(R) = 1.000$ ,  $B(S) = 500$
- 10 tuples per block, so  $T(R) = 10000$  and  $T(S) = 5.000$
- Assume  $V(S, B) = 100$
- Cost is  $1000 + \frac{10.000 \times 500}{100} = 50.000$
- Much higher than previous algorithms!
- Might be useful if  $R$  is much smaller than  $S$

How can we improve this

Use a sorted index, in which we get the tuples in order

Example

- Suppose that we have sorted indexes for  $B$  on both  $R$  and  $S$
- Let the  $B$ -values for the tuples in  $R$  be 1, 3, 4, 4, 5, 6
- Let the  $B$ -values for the tuples in  $S$  be 2, 2, 4, 4, 6, 7
- Find the first tuple in both relations
- Since  $1 < 2$  go on to the next tuple in  $R$
- Since  $3 > 2$ , go on to the next tuple in  $S$
- Since  $3 < 4$  go on to the next tuple in  $R$
- Retrieve tuples in  $R$  and  $S$  and join them

Note that all except the last part only use the index

Assuming that all tuples that match in the last step fit in memory, this costs basically  $T(R) + T(S)$

With a clustered index, we can read a block at a time, with cost  $B(R) + B(S)$  (plus a small cost for reading the index)

## Physical query plan

We outline remaining steps, namely

- Decisions whether to materialize intermediate results (store on disk) and when to pipeline to the next operator
- A notation for physical-query-plan operators, including access methods for relations, and algorithms for relational-algebra operators

## Pipelining vs. materialization

- Naive approach: implement each operation and store the result on disk until it is needed
- It is usually better to interleave the execution of several operations. The tuples produced by one operation are passed directly to the operation that uses it, without storing tuples on disk
- Pipelining is typically implemented by a network of iterators which call each other when needed
- This saves a lot of disk I/O
- Disadvantage: Less memory available for each operation



## Pipelining unary operations


- Projection and selection
- One block in main memory used for a block of data
- The next operator is implemented in such a way that `GetNext()` gets the first tuple in main memory, instead of accessing the disk

## Pipelining binary operations

- Use one block to pass result to next operator
- Number of blocks to compute and to use the result can vary, depending on size of result and sizes of arguments
- Example

$$(R(w, x) \bowtie S(x, y)) \bowtie U(y, z)$$

- $B(R) = 5.000, B(S) = B(U) = 10.000$
- $B(R \bowtie S) = k.$
- The algorithm depends on the value of  $k$
- Joins implemented as hash-joins
- $M = 100$



- $R \bowtie S$ . Neither relation fits in memory, so we use 2-pass hash-join
- First pass: Partition  $R$  into 100 buckets
- Each bucket takes 50 blocks
- Second pass of  $R \bowtie S$  uses 51 blocks:
  - 50 blocks for a bucket of  $R$
  - One block of  $S$
- 50 left to use for the join with  $U$

First case:  $k \leq 49$

- Assume  $k \leq 49$
- Pipeline them into 49 buffers, organize as hash table
- One block left for  $U$
- So second join can be evaluated in one pass
- Cost
  - 45.000 for the 2-pass hash-join of  $R$  and  $S$
  - 10.000 to read  $U$
- Total cost 55.000

Second case:  $49 < k < 5.000$

- Why 5.000?
- $5.000 = 100 \times 50$ , i.e.,  $M \times 50$
- This means that if we use a 50-bucket hash-join algorithm for joining  $(R \bowtie S)$  and  $U$ , we can fit an entire bucket of  $R \bowtie S$  in main memory
- The first step is to hash  $U$  into 50 buckets (each one takes 200 blocks)
- Now join  $R$  with  $S$  via two pass hash-join as before
- This uses 51 blocks, leaving 50 free. These blocks will be used to store the 50 buckets of  $R \bowtie S$
- Whenever we generate a tuple in  $R \bowtie S$ , put it in the appropriate main-memory block, writing blocks to disk when they are full

- Now join  $R \bowtie S$  with  $U$  bucket by bucket
  - Buckets of  $R \bowtie S$  take 100 blocks
  - Load one block of the corresponding bucket of  $U$  at each step
  - Total memory needed at each step: 101 blocks
- Costs
  - Read  $U$  and write tuples into buckets: 20.000
  - Two-pass hash-join of  $R$  and  $S$ : 45.000
  - Write out buckets of  $R \bowtie S$ :  $k$
  - Reading both relations in final join:  $k+10.000$
- Total  $75.000 + 2k$
- Big jump from 49 to 50 (but we could probably still use 1-pass algorithm with some thrashing)

Third case:  $k > 5.000$

- Note that we can't use a 2-pass algorithm in 50 buffers
- We could use a 3-pass one
- We haven't studied these. The cost would be  $20.000 + 2k$  I/O, but we'll see how to do better *without* pipelining
- Compute  $R \bowtie S$  using 2-pass hash-join and store on disk
- Cost  $45.000 + k$
- Join  $R \bowtie S$  with  $U$  using 2-pass hash-join with  $U$
- Since  $B(U) = 10.000$ , using 100 buckets we can fit a bucket in memory (we are using  $U$  as the *left* argument)
- This join then costs  $30.000 + 3k$  (read each argument twice, write once)
- Total cost  $75.000 + 4k$

## Summary

Range	Pipeline?	Final join	I/O
$k \leq 49$	Pipeline	1-pass	55.000
$50 \leq k \leq 5.000$	Pipeline	50-bucket 2-pass	$75.000 + k$
$5.000 < k$	Materialize	100-bucket 2-pass	$75.000 + 4k$



## Physical Query Plans: Notation

- Each operator becomes an operator of the physical plan
- Each leaf becomes a scan operator applied to a relation
- Materialized relation indicated by Store operator, plus the appropriate scan operator (TableScan unless we have constructed an index)
- We skip the latter detail, and indicate materialization by double lines ( “crossing out” ); all others are pipelined

## Operations for leaves

Possible operations on leaves (base relations)

- $\text{TableScan}(R)$ : Read all blocks in arbitrary order
- $\text{SortScan}(R, L)$ : Tuples read in sorted order, using the attributes in  $L$
- $\text{IndexScan}(R, C)$ :  $C$  is a condition of the form  $A < c$ ,  $A = c$  or  $A > c$ . Tuples are accessed using an index on  $A$
- $\text{IndexScan}(R, A)$ : The entire relation  $R$  is retrieved via an index on  $R.A$  (identical to  $\text{TableScan}$ , as far as we are concerned)

## Selection


- $\sigma_C(R)$  can be combined with the access method
- If the argument is not a stored relation, or we have no index, we use the `Filter` operator
- Conjunction, e.g.  $A\theta_C$  AND  $D$  with index on  $R.A$ :
  - Use `IndexScan(R,C)` to access  $R$  with first condition followed by
  - `Filter(D)`

## Physical Sort Operators


- `SortScan(R,L)`
- Sort-based join or grouping: use an explicit `Sort(L)`
- This can also be used for `ORDER BY` clause

Other operations. Specify

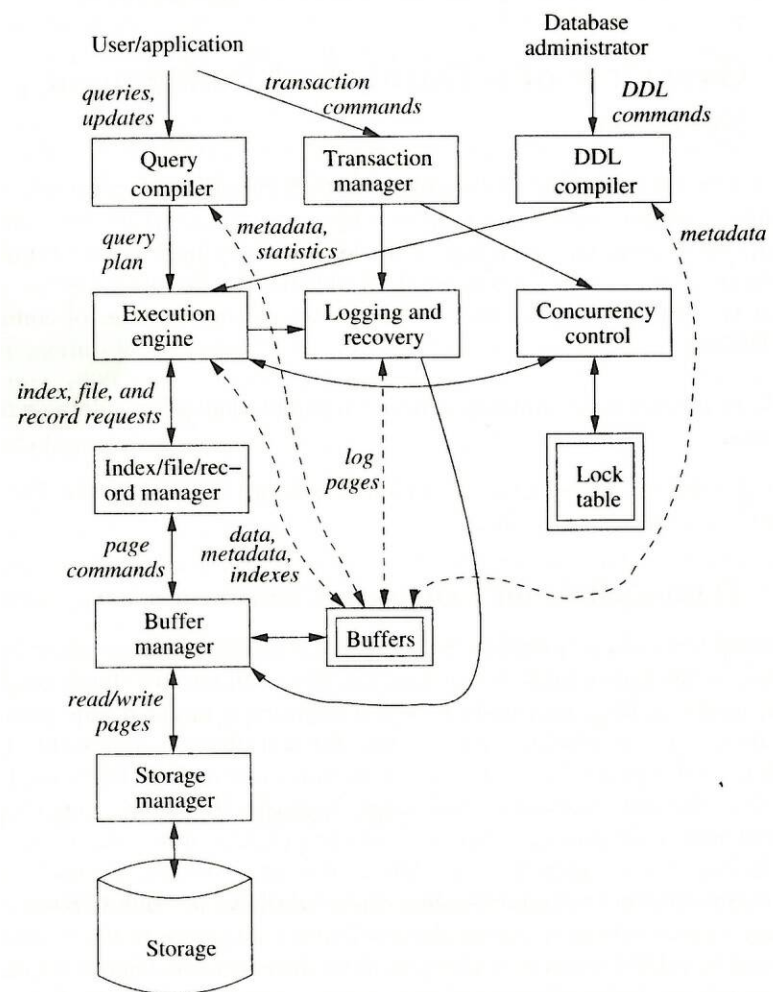
- The operation being performed (join, grouping, etc.)
- Parameters (condition in theta-join etc.)
- Strategy (sort-based etc.)
- Number of passes (can be left until run-time)
- Expected number of memory blocks

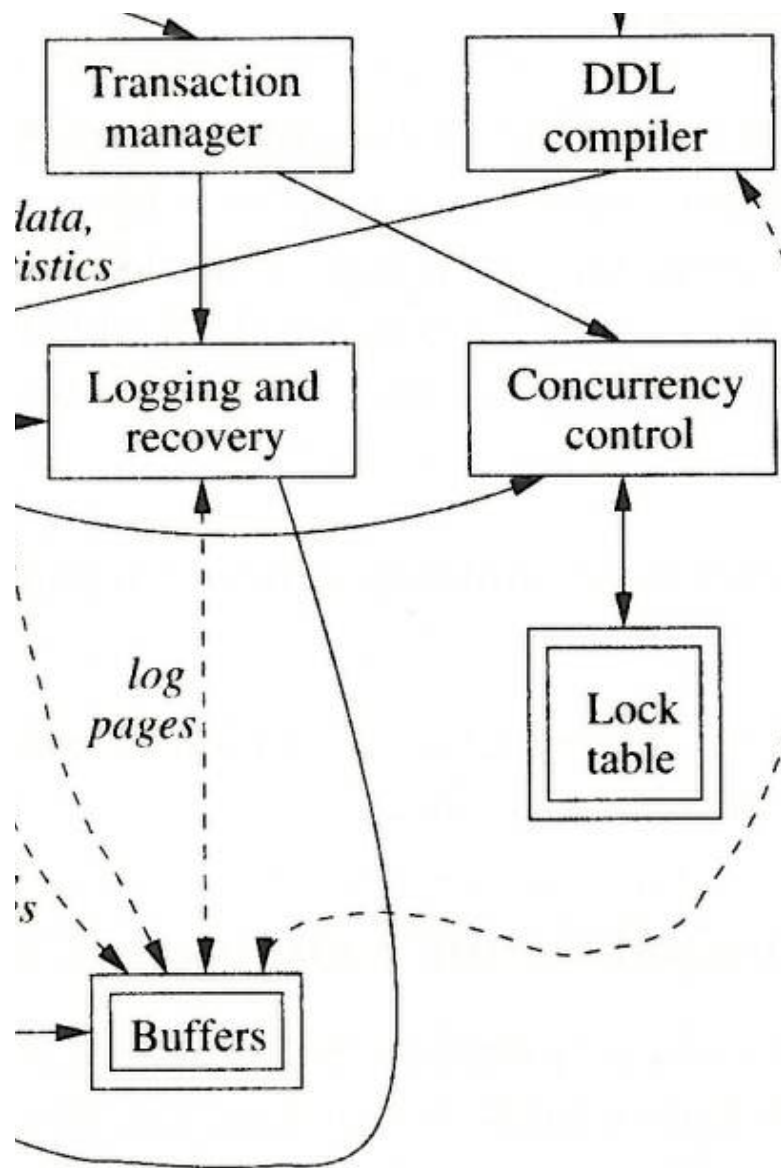


- Figure shows our example when  $k > 5.000$
- Access each of the relations by a table-scan
- Use two-pass hash-join for the first join and materialize it (crossed lines)
- Use two-pass hash-join for the second join (materialization implies table-scan)



- The case  $k < 49$
- Second join has a different number of passes, a different number of memory blocks
- Left argument is pipelined







## Concurrency control

- Our model so far: one user queries or modifies the database
- Operations on the database are executed one at a time, and the database state left after an operation is the state upon which the next acts
- Operations are “atomic”: It is impossible for the system to fail in the middle of a modification, leaving the database in an inconsistent state
- Real life is different
- In applications like Web services, banking, or airline reservations, hundreds of operations per second may be performed on the database
- Two operations affecting the same bank account or flight may be executed at the same time, and may might interact in strange ways

## Example

- Example of what could go wrong
- Note that we combine the DBMS with the O/S to get an intuitive example (the DBMS really only deals with SQL statements, but similar problems arise even there)
- Airline website where customers chose a seat
- Relation

Flights(fltNo, fltDate, seatNo, seatStatus)

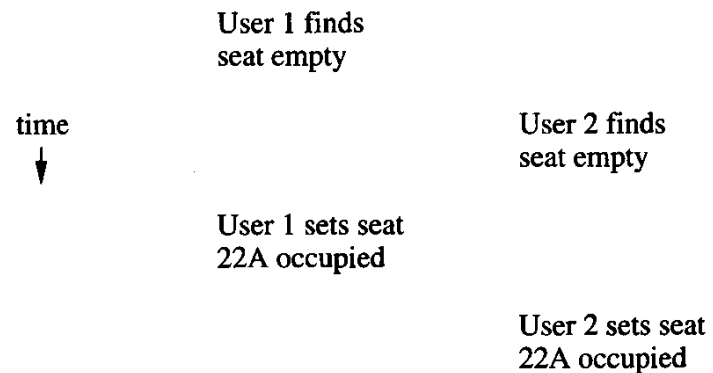
- Query

```
SELECT seatNo
FROM Flights
WHERE fltNo = 123 AND fltDate = DATE 2008-12-25
      AND seatStatus = available;
```

- Customer clicks on empty seat, say 22A
- System reserves this seat, and modifies the database

```
UPDATE Flights
SET seatStatus = occupied
WHERE fltNo = 123 AND fltDate = DATE 2008-12-25
    AND seatNo = 22A;
```

- Suppose another customer is using the system and the same time
- He also sees seat 22A empty and clicks on it



- Both customers believe they have got seat 22A
- Only the second one “really” has it
- In SQL: *transaction*, a group of operations that must be performed together
- If these are defined as transactions, the two transactions must behave as though they were run *serially*
- But they don’t have to be actually run serially; we can allow overlap if it doesn’t create problems
- The more overlap, the better the performance of the system
- A sequence of operations (as in our example) is *serializable* if it is equal to a serial execution of the transactions

## Another example

- Bank account records

Accounts(acctNo, balance)

- Transferring \$100 from account 123 to account 456
- First check whether there is at least \$100 in account 123
- If so, we execute the following two steps:

- Add \$100 to account 456:

```
UPDATE Accounts  
SET balance = balance + 100  
WHERE acctNo = 456;
```

- Subtract \$100 from account 123:

```
UPDATE Accounts  
SET balance = balance - 100  
WHERE acctNo = 123;
```

## Possible problem

- Suppose there is a failure between the two steps
- The database is left in a state where money has been transferred into the second account, but taken out of the first account
- The two updates should be done atomically: either both are done or neither

## Transactions

- Solution: group database operations into *transactions*
- A transaction is a collection of one or more operations on the database that must be executed atomically: either all operations are performed or none are
- SQL (default) requires that transactions are executed in a serializable manner: equivalent to one transaction executed in its entirety before the other
- Note that this does not specify *which* transaction is executed first
- The DBMS may allow the user to specify less stringent constraints on the interleaving of operations from two or more transactions

## Outline

We study

- Theory of serializability: how to test
- Specification in SQL
- Theory of locking: how to avoid nonserializable behaviour



## Testing serializability

- Schedule: An execution of transactions with operations interleaved

$T_1$	$T_2$
READ( $A, t$ )	READ( $A, s$ )
$t \leftarrow t + 100$	$s \leftarrow s * 2$
WRITE( $A, t$ )	WRITE( $A, s$ )
READ( $B, t$ )	READ( $B, s$ )
$t \leftarrow t + 100$	$s \leftarrow s * 2$
READ( $B, t$ )	READ( $B, s$ )

- Column 1 is transaction  $T_1$ , column 2  $T_2$
- Nothing is implied about execution time
- Invariant: If  $A = B$  before executing a transaction, it also holds afterwards
- The invariant is not part of the theory: it is only used as an easy way to see when an execution is wrong

## Serial schedules

- Our tables now describe which operation is executed at each step, from top to bottom
- We also list the values of  $A$  and  $B$  at each step
- There are two serial executions:  $T_1$  followed by  $T_2$  and  $T_2$  followed by  $T_1$

# Serial schedule

$T_1$	$T_2$	A	B
		25	25
READ( $A, t$ )			
$t \leftarrow t + 100$			
WRITE( $A, t$ )		125	
READ( $B, t$ )			
$t \leftarrow t + 100$			
READ( $B, t$ )			125
	READ( $A, s$ )		
	$s \leftarrow s * 2$		
	WRITE( $A, s$ )	250	
	READ( $B, s$ )		
	$s \leftarrow s * 2$		
	READ( $B, s$ )		250

## Another serial schedule

Note that the final results are different

$T_1$	$T_2$	A	B
		25	25
	READ( $A, s$ )		
	$s \leftarrow s * 2$		
	WRITE( $A, s$ )	50	
	READ( $B, s$ )		
	$s \leftarrow s * 2$		
	READ( $B, s$ )		50
READ( $A, t$ )			
$t \leftarrow t + 100$			
WRITE( $A, t$ )		150	
READ( $B, t$ )			
$t \leftarrow t + 100$			
READ( $B, t$ )			150

## Serialiazble schedule

$T_1$	$T_2$	A	B
		25	25
READ( $A, t$ )			
$t \leftarrow t + 100$			
WRITE( $A, t$ )		125	
	READ( $A, s$ )		
	$s \leftarrow s * 2$		
	WRITE( $A, s$ )	250	
READ( $B, t$ )			
$t \leftarrow t + 100$			
READ( $B, t$ )			125
	READ( $B, s$ )		
	$s \leftarrow s * 2$		
	READ( $B, s$ )		250

- $A = B$  holds
- This doesn't prove serializability, but we shall see later that it is indeed equivalent to  $(T_1, T_2)$

# Non-serializable schedule

$T_1$	$T_2$	A	B
		25	25
READ( $A, t$ )			
$t \leftarrow t + 100$			
WRITE( $A, t$ )		125	
	READ( $A, s$ )		
	$s \leftarrow s * 2$		
	WRITE( $A, s$ )	250	
	READ( $B, s$ )		
	$s \leftarrow s * 2$		
	READ( $B, s$ )		50
READ( $B, t$ )			
$t \leftarrow t + 100$			
READ( $B, t$ )			150

$A \neq B!$

# A slightly different example

$T_1$	$T_2$	A	B
		25	25
READ( $A, t$ )			
$t \leftarrow t + 100$			
WRITE( $A, t$ )		125	
	READ( $A, s$ )		
	$s \leftarrow s + 200$		
	WRITE( $A, s$ )	325	
	READ( $B, s$ )		
	$s \leftarrow s + 200$		
	READ( $B, s$ )		225
READ( $B, t$ )			
$t \leftarrow t + 100$			
READ( $B, t$ )			325

It can be shown that this is equivalent to  $(T_1, T_2)$

## Testing

- Analysing the details of the arithmetic is very difficult
- It is quite rare to have a situation like the previous one
- We therefore ignore the details, and test if the schedule will be serializable, no matter what the program does
- We analyze *only* the sequence of reads and write
- This means we miss schedules like the previous one



## General principle

- Our algorithms must not give *false positives*: Saying a schedule is serializable when it is not
- A few *false negatives* are OK. Saying a schedule is not serializable when it simply reduces performance, but never gives false behaviour
- So we seek to design algorithms with no false positives, and with a few false negatives as we can
- Analyzing only the sequences of reads and writes is our first example of this rule

## Notation

- $r_T(X)$ : Transaction  $T$  reads database element  $X$
- $w_T(X)$ : Transaction  $T$  writes database element  $X$
- If the transactions are  $T_1, T_2, \dots$ , we write  $w_i(X)$  instead of  $w_{T_i}(X)$
- Example of transactions.

- $T_1$ :

$$r_1(A); w_1(A); r_1(B); w_1(B)$$

- $T_2$ :

$$r_2(A); w_2(A); r_2(B); w_2(B)$$

- Schedule (serializable):

$$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$$

## Notation

- An action is an expression of the form  $r_i(X)$  or  $w_i(X)$
- A transaction  $T_i$  is a sequence of actions with subscript  $i$
- A schedule  $S$  of a set of transactions  $T$  is a sequence of actions, in which, for each transaction  $T_i$  in  $T$ , the actions of  $T_i$  appear in  $S$  in the same order that they appear in the definition of  $T_i$  itself
- We say that  $S$  is an interleaving of the actions of the transactions of which it is composed.

## Conflict-serializability

Given a schedule  $S$ , examine *consecutive* actions from different transactions  $T_i$  and  $T_j$  ( $i \neq j$ )

- $r_i(X); r_j(Y)$  is never a conflict, since neither changes the value of any database element
- $r_i(X); w_j(Y)$ ,  $X \neq Y$  is not a conflict:  $T_j$  writes  $Y$  before  $T_i$  reads  $X$ , the value of  $X$  is not changed. Furthermore the read of  $X$  by  $T_i$  has no effect on  $T_j$
- $w_i(X); r_j(Y)$ ,  $X \neq Y$  is not a conflict for the same reason
- $w_i(X); w_j(Y)$ ,  $X \neq Y$  is not a conflict for similar reasons

If two actions do not conflict, we can swap them without changing the results

## Conflicts

- Two actions of the same transaction, such as  $r_i(X); w_i(Y)$  always conflict. We cannot change the order of actions in a single transaction
- $w_i(X); w_j(X)$  is a conflict. The value of  $X$  is whatever  $T_j$  computed. If we swap the order, we leave  $X$  with the value computed by  $T_i$
- $r_i(X); w_j(X)$  is a conflict. If we swap the order, the value of  $X$  read by  $T_i$  will be that written by  $T_j$ , which may be different from the previous value of  $X$
- $w_i(X); r_j(X)$  is a conflict, for similar reasons

- Any two actions of different transactions may be swapped unless:
  - They involve the same database element, and
  - At least one is a write.
- Given schedule, we make as many nonconflicting swaps as needed, with the goal of turning the schedule into a serial schedule.
- If we can do so, then the original schedule is serializable
- Two schedules are *conflict-equivalent* if each can be turned into the other by a sequence of nonconflicting swaps of adjacent actions
- A schedule is *conflict-serializable* if it is conflict-equivalent to a serial schedule
- Conflict-serializability implies serializability
- The converse is false, but conflict-serializability is usually sufficient in practice (remember that a few false negatives are not a problem)

Converse is false

$T_1$ ,  $T_2$ , and  $T_3$

- $T_1$ :  $w_1(Y); w_1(X)$
- $T_2$ :  $w_2(Y); w_2(X)$
- $T_3$ :  $w_1(X)$
- $S_1$ :  $w_1(Y); w_1(X); w_2(Y); w_2(X); w_3(X)$  is serial
- $S_2$ :  $w_1(Y); w_2(Y); w_2(X); w_1(X); w_3(X)$  is not serial
- $S_2$  is not conflict-serializable
- But the two are equivalent, as the final effect is for  $X$  to get the value written by  $T_3$ ;  $T_1$  and  $T_2$  are completely irrelevant

### Example

$S: r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$

$S$  is conflict serializable, and equivalent to  $(T_1, T_2)$

### Proof

$r_1(A);$	$w_1(A);$	$r_2(A);$	<u><math>w_2(A);</math></u>	<u><math>r_1(B);</math></u>	$w_1(B);$	$r_2(B);$	$w_2(B)$
$r_1(A);$	$w_1(A);$	<u><math>r_2(A);</math></u>	<u><math>r_1(B);</math></u>	$w_2(A);$	$w_1(B);$	$r_2(B);$	$w_2(B)$
$r_1(A);$	$w_1(A);$	$r_1(B);$	<u><math>r_2(A);</math></u>	<u><math>w_1(B);</math></u>	$w_2(A);$	$r_2(B);$	$w_2(B)$
$r_1(A);$	$w_1(A);$	$r_1(B);$	$w_1(B);$	$r_2(A);$	$w_2(A);$	$r_2(B);$	$w_2(B)$



## Testing conflict-serializability

Given schedule  $S$ , involving transactions  $T_1$  and  $T_2$  (and maybe others)  $T_1$  *takes precedence over*  $T_2$ , written  $T_1 <_S T_2$ , if there are actions  $A_1$  of  $T_1$  and  $A_2$  of  $T_2$  such that

- $A_1$  is ahead of  $A_2$  in  $S$
- Both  $A_1$  and  $A_2$  involve the same database element
- At least one of  $A_1$  and  $A_2$  is a write action.

Note that these are exactly the conditions under which we cannot swap the order of  $A_1$  and  $A_2$

Therefore  $A_1$  appears before  $A_2$  in any schedule that is conflict-equivalent to  $S$

So any conflict-equivalent serial schedule must have  $T_1$  before  $T_2$

## Precedence graph

A directed graph

- Nodes: transactions of  $S$
- Edge from node  $T_i$  (written  $i$ ) to  $T_j$  when  $T_i <_S T_j$

Example:

$S$ :  $r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

- $T_2 <_S T_3$  due to  $r_2(A)$  being ahead of  $w_3(A)$  (and other reasons)
- $T_1 <_S T_2$  due to  $r_1(B)$  being ahead of  $w_2(B)$

Graph



## Algorithm

- Construct the precedence graph
- Test if it has (directed) cycles
- If it does,  $S$  is not conflict serializable
- If not, any topological order of the nodes is a conflict-equivalent serial order




- Our previous graph is acyclic, so the schedule is serializable
- Only one order possible:  $(T_1, T_2, T_3)$
- We can convert the schedule:  
 $S: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$
- Into  
 $S': r_1(B); w_1(B); r_2(A); r_2(B); w_2(B); w_2(A); r_3(A); w_3(A)$
- By a series of local swaps

## Another example

- Consider the schedule

$S: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$

- $A: T_2 \prec_S T_3$
- $B: T_1 \prec_S T_2, T_2 \prec_S T_1$
- Graph is cyclic, so schedule is not conflict-serializable



## Why this works

First direction: If there is a cycle, the schedule is not serializable

- Suppose there is a cycle of length  $n$ ,  $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n \rightarrow T_1$
- Then the actions of  $T_1$  must precede those of  $T_2$ , which must precede those of  $T_3$  etc.
- In this way, the actions of  $T_1$  must precede those of  $T_1$ , a contradiction

Converse: Proof by induction on the number of actions

- Induction: If the graph for  $n$  has no cycles, we can reorder the schedules actions using legal swaps so the schedule becomes serial
- Basis:  $n = 1$ . Schedule must already be serial
- Induction. Assume the claim true for  $n = 1$ . We prove it for  $n$
- Schedule consists of actions of transactions  $T_1, \dots, T_n$  with acyclic graph  $S$
- If the graph is acyclic, it must have a node  $i$  with no incoming arcs
- No incoming arcs means that no node precedes  $i$  in our ordering

- But this means that there is no action involving another transaction  $T_j$  that
  - Precedes some action of  $T_i$  and
  - conflicts with that action
- We can then move all the actions of  $T_1$  to the beginning of the schedule, getting

(Actions of  $T_i$ )(Actions of the remaining  $n - 1$  transactions)

- Consider the second part
- The relative order of actions is preserved, so the graph is the original graph with  $i$  deleted
- This graph must be acyclic, so we can apply the inductive hypothesis
- This yields a serial ordering