
The Entity Relationship Model



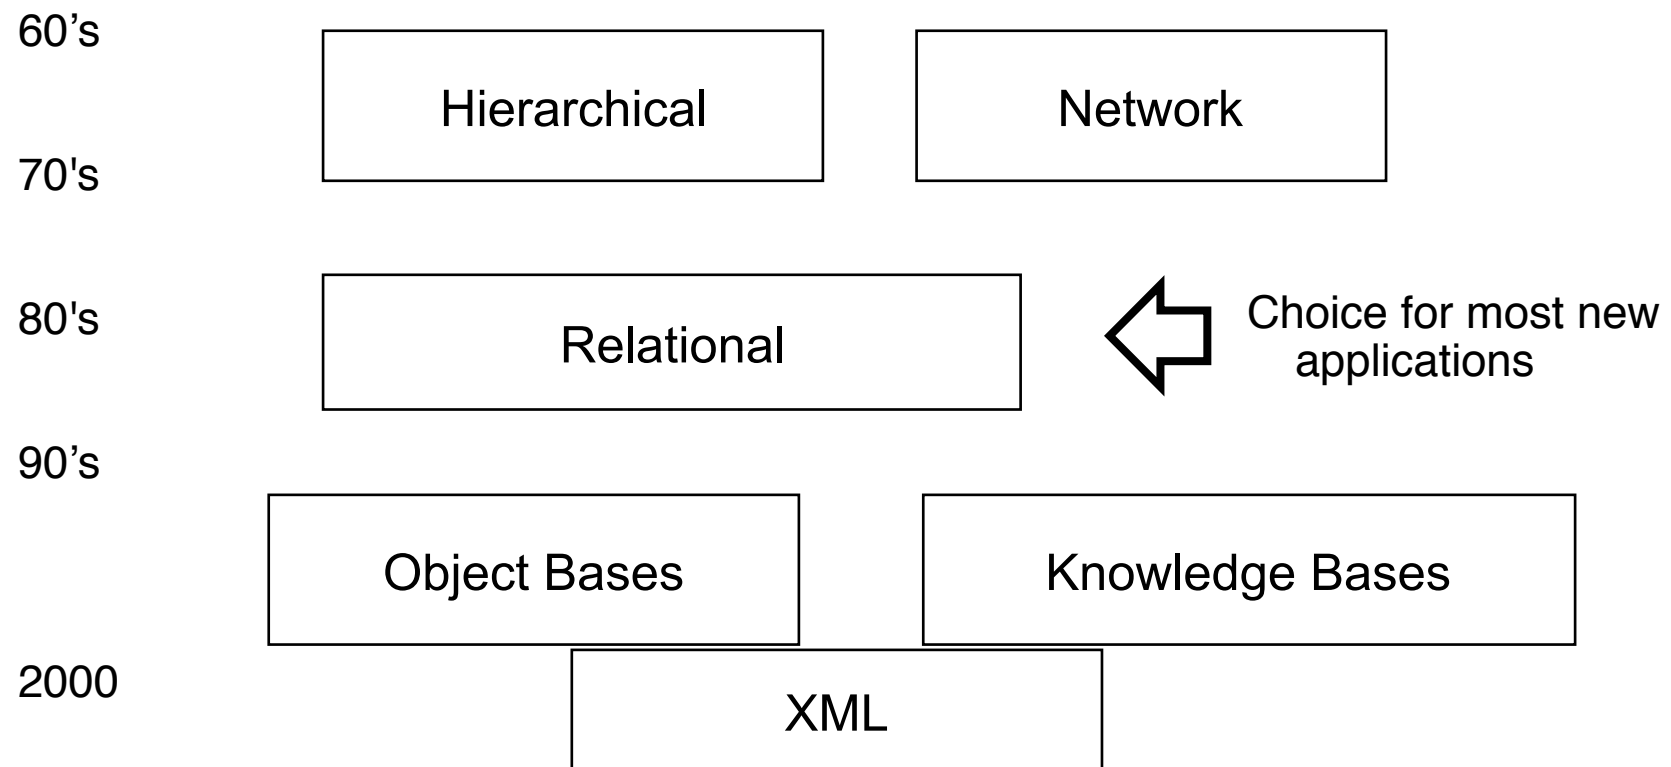
RDBMS

What is a Database Management System?

1. Manages **very large** amounts of **data**.
2. Supports **efficient access** to very large amounts of data.
3. Supports **concurrent access** to very large amounts of data.
Example: bank and its ATM machines.
4. Supports **secure, atomic access** to very large amounts of data.
Example: Contrast two people editing the same UNIX file – last to write “wins” – with the problem if two people deduct money from the same account via ATM machines at the same time – new balance is wrong whichever writes last.


Data Models: Historic Overview

A data model says **what** information is to be contained in a database, **how** the information will be used, and how the items in the database will be **related** to each other



Now?

Why Study Data Management??

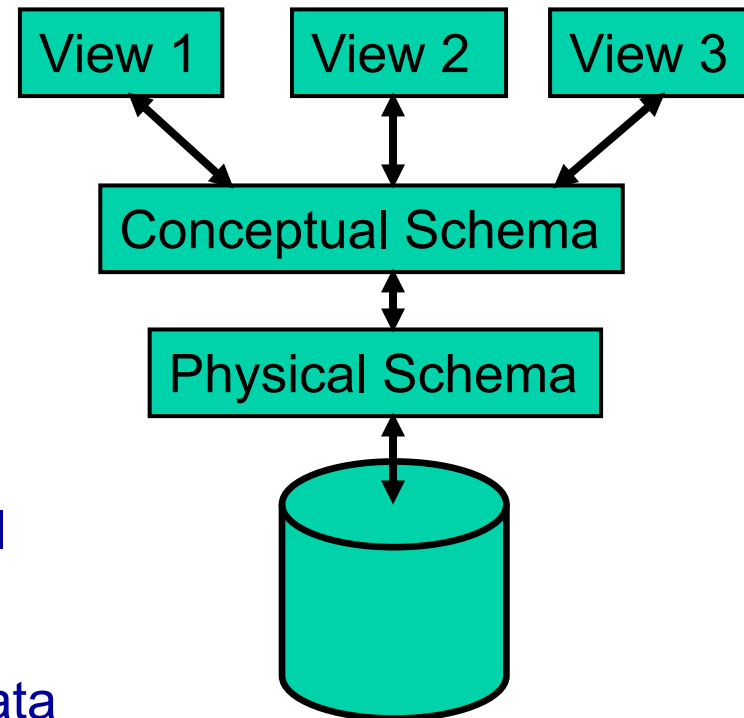


- Shift from computation to information
 - ◆ at the “low end”: scramble to webspace (a mess!)
 - ◆ at the “high end”: scientific applications
- Datasets increasing in diversity and volume.
 - ◆ Digital libraries, interactive video, Human Genome project, EOS project
 - ◆ ... need for DBMS exploding
- DBMS encompasses most of CS
 - ◆ OS, languages, theory, AI, multimedia, logic

Levels of Abstraction

Many views,
single conceptual schema
and physical schema.

- Views describe how users see the data.
- Conceptual schema defines logical structure
- Physical schema describes how data are stored



Example: University Database

- External Schema (View):
 - ♦ *Course_info(cid:string,enrollment:integer)*
- Conceptual schema:
 - ♦ *Students(sid: string, name: string, login: string, age: integer, gpa:real)*
 - ♦ *Courses(cid: string, cname:string, credits:integer)*
 - ♦ *Enrolled(sid:string, cid:string, grade:string)*
- Physical schema:
 - ♦ Relations stored as unordered files.
 - ♦ Index on student id (sid).

Relational Model Example

Relational model is based on **tables**...

e.g. *CustomerTable* in a Bank Database

acct #	name	balance
12345	Tasos	\$1000.2
34567	Yannis	\$285.48
...

Today used in *most* DBMS's...

The DBMS Marketplace

- Relational DBMS companies – Oracle, Sybase – are among the largest software companies in the world.
- IBM offers its relational DB2 system. With IMS, a nonrelational system, IBM is by some accounts the largest DBMS vendor in the world.
- Microsoft offers SQL-Server, plus Microsoft Access for the cheap DBMS on the desktop, answered by “lite” systems from other competitors.
- Relational companies also challenged by “object-oriented DB” companies.
- But countered with “object-relational” systems, which retain the relational core while allowing type extension as in OO systems.

Structured Query Language (SQL)

acct #	name	balance
12345	Tasos	\$1000.2
34567	Yannis	\$285.48
...

SELECT *name*
FROM *CustomerTable*
WHERE *balance* >= 500

SELECT *acct#*
FROM *CustomerTable*
WHERE *name* = "Yannis" **AND**
balance >= 200

Three Aspects to Studying DBMS's


- Modeling and design of databases Management System.
 - ◆ Allows exploration of issues before committing to an implementation.
- Programming: queries and DB operations like update.
 - ◆ SQL = “intergalactic dataspeak.”
 - ◆ Is it enough?
- Data Management System Implementation.
 - ◆ Is a database enough? Or we need something more?

The Entity Relationship Model

Entity/Relationship Model


Main idea: Diagrams to represent designs.

- *Entity* like object, = “thing”
- *Entity set* like class = set of “similar” entities/objects.
- *Attribute* = property of entities in an entity set, similar to fields of a struct.
- In diagrams, entity set → rectangle; attribute → oval.



Relationships

- Connect two or more entity sets.
- Represented by diamonds.




Relationship Set

- Think of the “value” of a relationship set as a table.
- One column for each of the connected entity sets.
- One row for each list of entities, one from each set, that are connected by the relationship.
- Example: stars-in relationship set


Movies	Stars
Basic Instinct	Sharon Stone
Total Recall	Arnold Schwarzenegger
Total Recall	Sharon Stone
...	...

Why are relationship sets important?


Multiplicity of Relationships



Many-many



Many-one




One-one

Representation of Many-One:

Arrow pointing to “one.”


Rounded arrow = “exactly one.”

... or simply



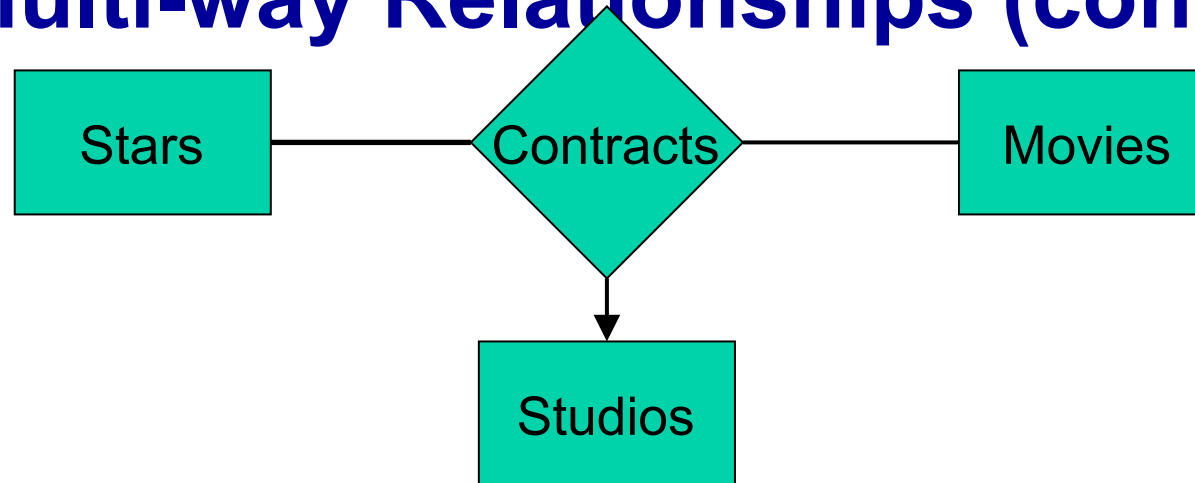
Multi-way Relationships

Usually binary relationships suffice.



- However, there are some cases where three or more E.S. must be connected by one relationship.
- Example: *Stars* have *Contracts* with *Studios* for particular *Movies*.

Multi-way Relationships (cont.)



Stars	Movies	Studios
Sharon Stone	Basic Instinct	Sony
Arnold Schwarzenegger	Total Recall	Columbia
Sharon Stone	Total Recall	Columbia
...

Multi-way Relationships (cont.)


Notice arrow convention for multi-way relationships:

“all other E.S. determine one of these.”

- Not sufficiently general to express any possibility.
- Assume that a studio, say, depended only on the movie, then we could use two 2-way relationships: *Studios-Movies* and *Movies-Stars*.
- Or better: just make *studio* an attribute of *Movies*.


Roles in relationships

Sometimes an E.S. participates more than once in a relationship. Thus, we need to label edges with *roles* to distinguish.



Husband	Wife
Arnold Schw	Sharon Stone
Arnold Schw	Jamie Lee Curtis
...	...

Roles in relationships (cont.)



Actor ₁	Actor ₂
Tom Cruise	Nicole Kidman
Nicole Kidman	Tom Cruise
...	...


Note: *Co-star* is symmetric, *Acted-as* was not.

There is no way to say “symmetric” in E/R.


Design Question:

What if we replace *Husband* and *Wife* by one role, e.g., *couple*?

Attributes on Relationships



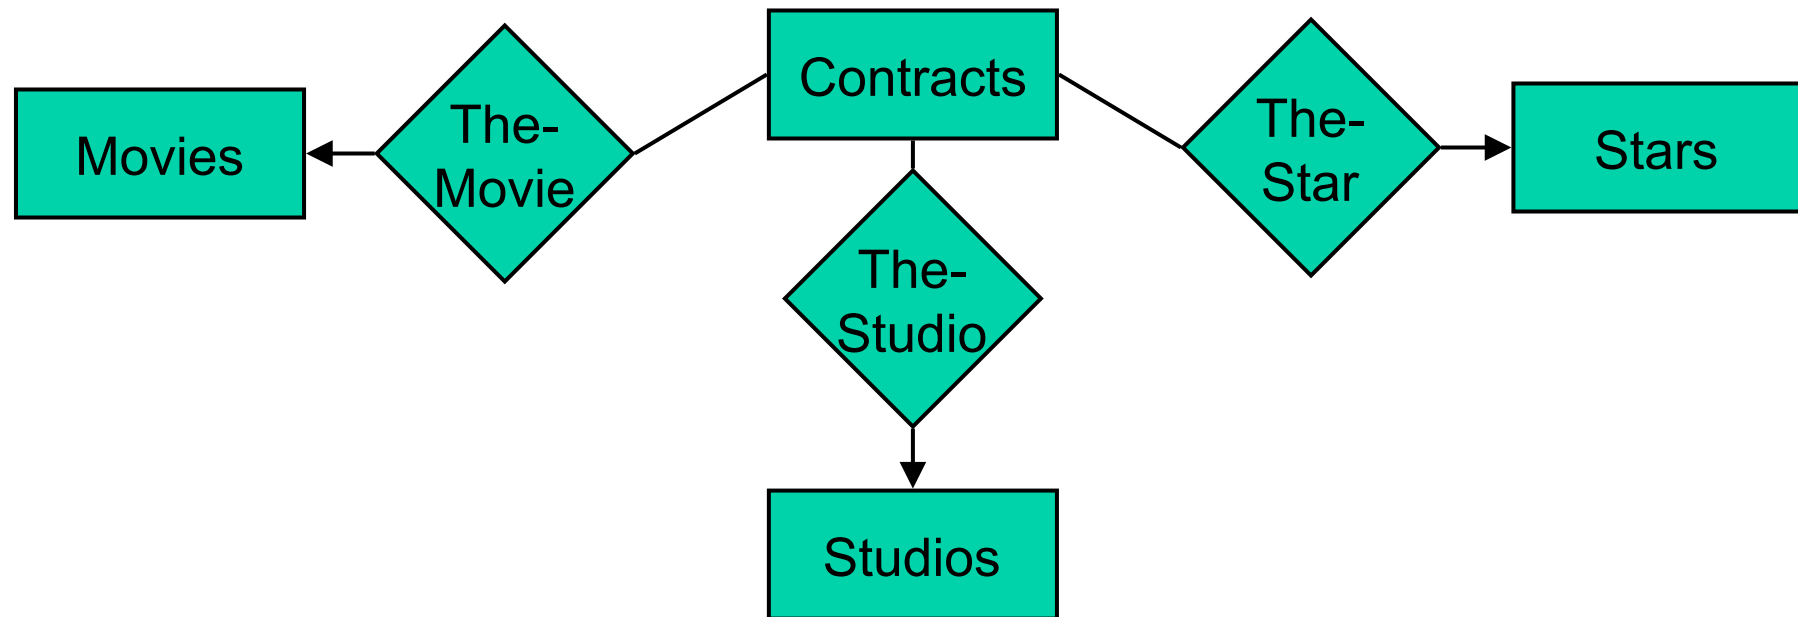
... a shorthand for the 3-way relationship:



Converting Multi-way to 2-Way

Baroque in E/R, but necessary in certain “object-oriented” models.

- Create a new E.S. to represent rows of a relationship set.
- Add many-one relationships from the new E.S. to the E.S.’s that participated in the original relationship.



More Design Issues

1. Subclasses.
2. Keys.
3. Weak entity sets.

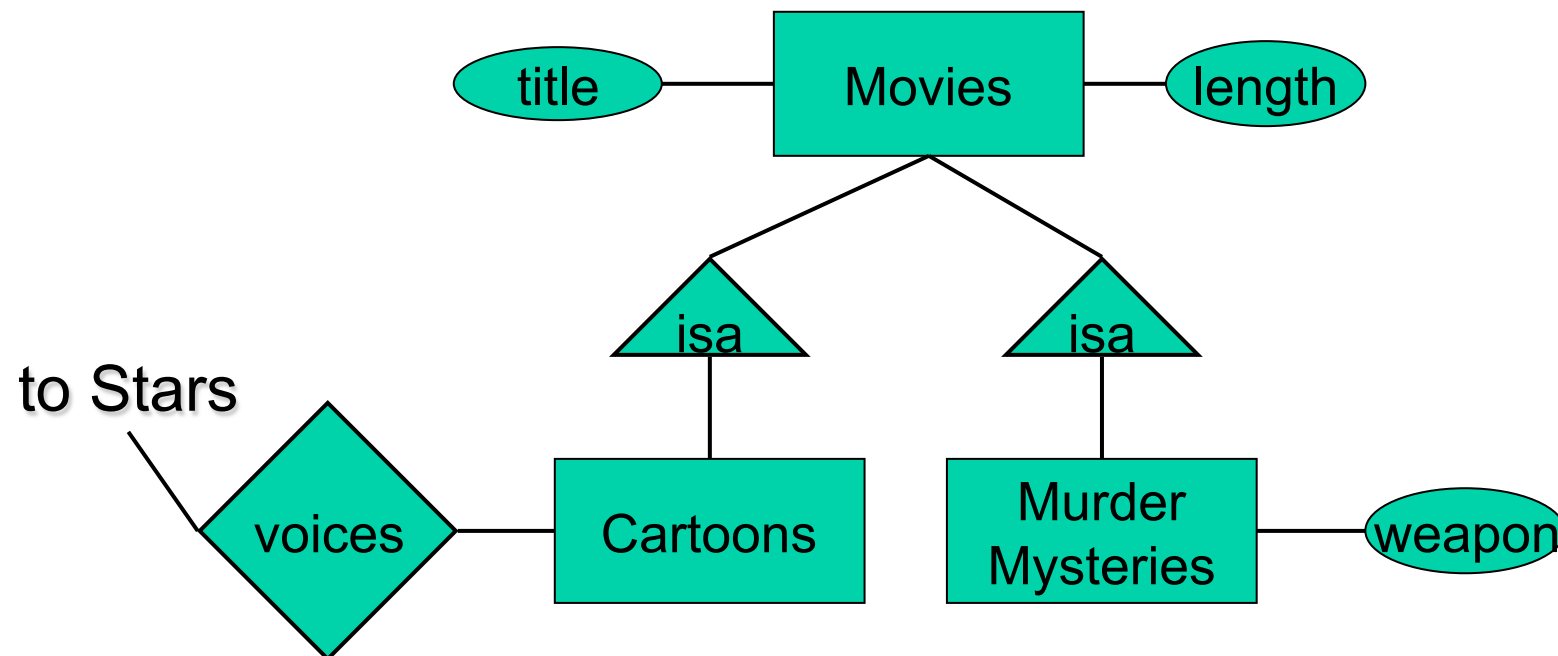
Subclasses

Subclass = special case = fewer entities =
more properties.

- Example: A cartoon is a kind of movie. In addition to the properties (= attributes and relationships) of movies, there is a *voices* attribute for cartoons.

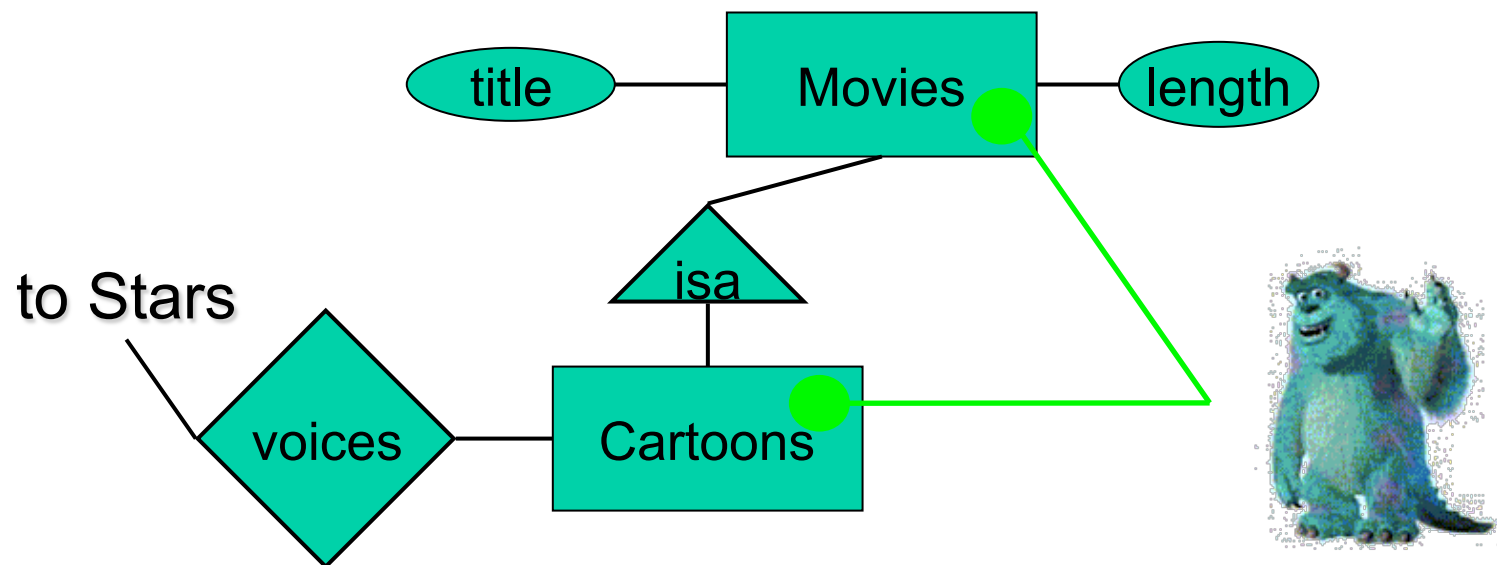
E/R Subclasses

- Subclasses form a tree (no multiple inheritance).
- *isa* relationships indicate the subclass relation
- they are one-one



Different Subclass Viewpoints

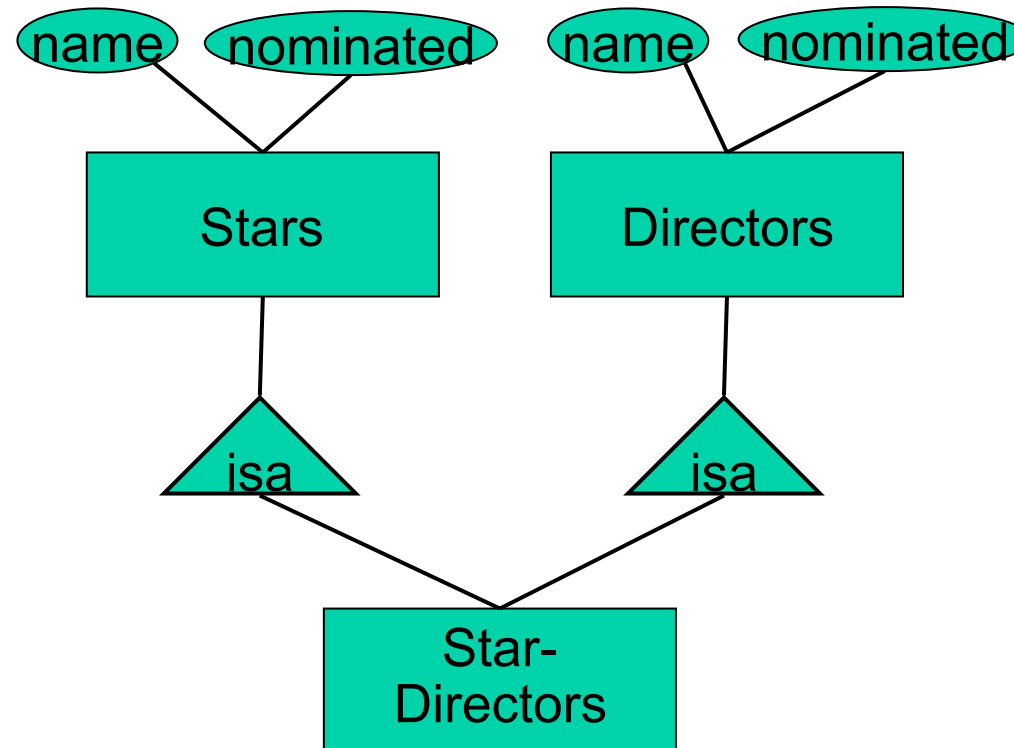
- *E/R viewpoint*: An entity has a *component* in each entity set to which it logically belongs. Its properties are the union of the properties of these E.S.
- *Object-oriented viewpoint*: An object (entity) belongs to exactly one class. It *inherits* properties of its superclasses.



Courtesy of Pixar

Multiple Inheritance

Theoretically, an E.S. could be a subclass of several other entity sets.



Problems

How should conflicts be resolved?

- Example: nominated means Oscar nominated as an actor for *Stars*, and as a director for *Directors*. What does it mean for *Star-Directors*?
- Need ad-hoc notation to resolve meanings.
- In practice, we shall assume a tree of entity sets connected by *isa*, with all “isa’s” pointing from child to parent.

Keys

Consider an entity set E . A *key* is a set of attributes K of E such that, given two entities e_1 and e_2 of E , e_1 and e_2 cannot have identical values in K

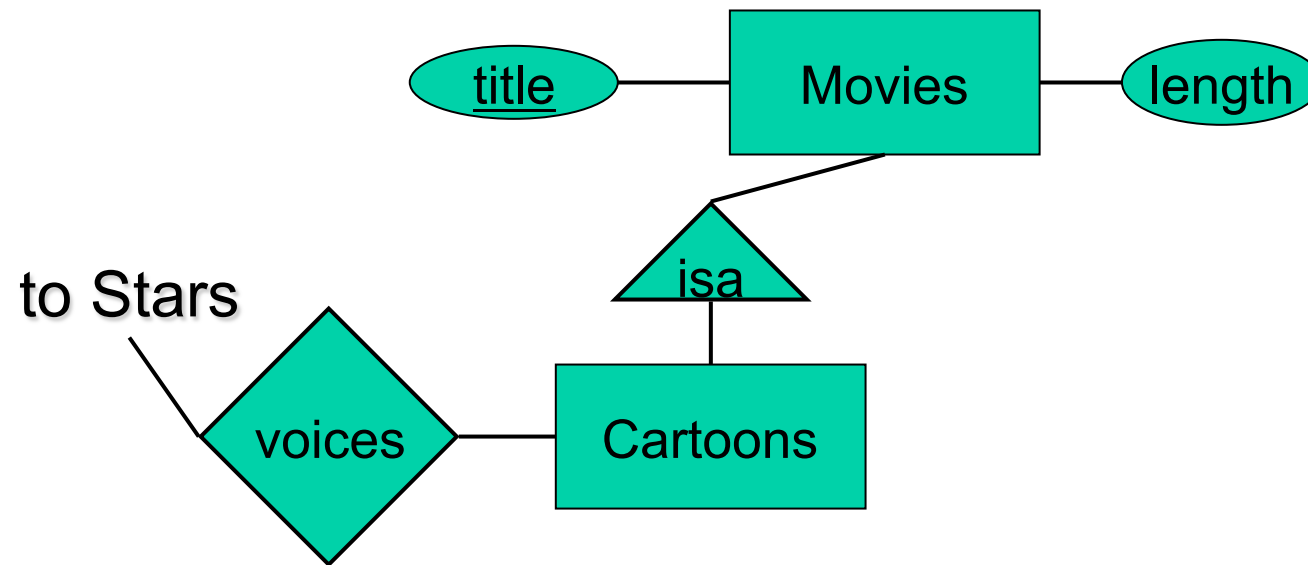
- In E/R model, every E.S. must have a key.
 - ◆ It could have more than one key, but one set of attributes is the “designated” key.
- In E/R diagrams, you should underline all attributes of the designated key.

Example


Suppose *title* is the key for *Movies*.

It makes sense to use *title* as key for *Cartoons* also.

Thus, we use root key as key for all subclasses.



Example: A Multi-attribute Key




Possibly, the combination of *title* + *length* also forms a key, but we have not designated it as such.

Example: Logins (Email Addresses)

Login name = user name + host name, e.g. velgias@disi.unitn.eu

- A “login” entity corresponds to a user name on a particular host. The *Login* entity doesn’t record the host, just the user name, e.g., tasos.
- Key for a login = the user name at the host (which is unique for that host only) + the IP address of the host (which is unique globally).



- Design issue: Under what circumstances could we simply make login-name and host-name be attributes of logins, and dispense with the **weak** E.S.?


Weak Entity Sets

Consider an entity set E . It's possible that E 's key comes not (completely) from its own attributes, but also from the keys of one or more E.S.'s to which E is linked (through a *supporting* many-one relationship).

- Entity set E is called a *weak* E.S.
- It is represented by putting double rectangle around E and a double diamond around each supporting relationship.
- Many-one-ness of supporting relationship (includes 1-1) essential.
 - ◆ With many-many, we wouldn't know which entity provided the key value.
- “Exactly one” also essential, or else we might not be able to extract key attributes by following the supporting relationship.

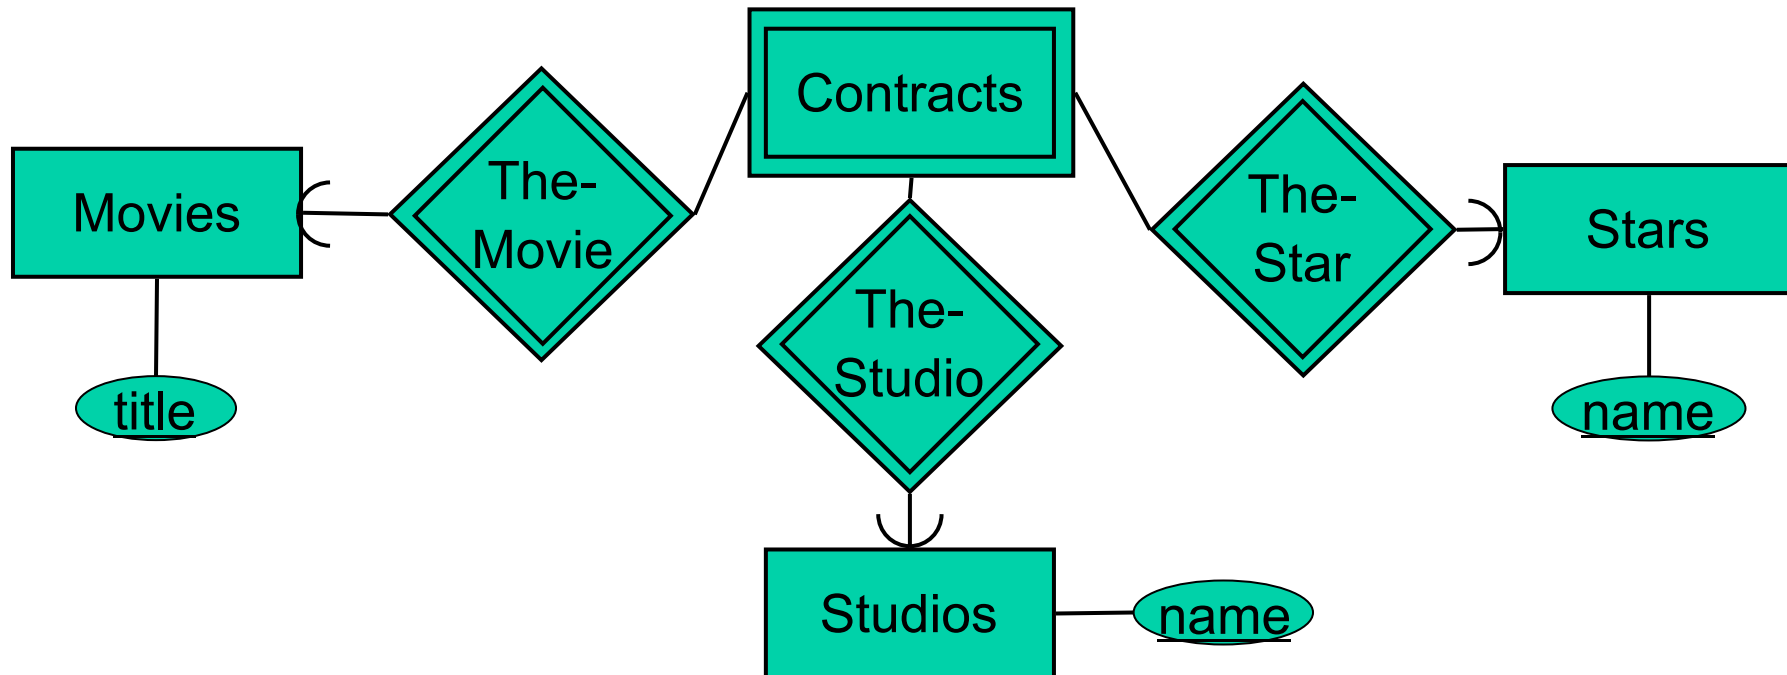
Example: Chain of “Weakness”

Consider IP addresses consisting of a primary domain (e.g., edu), subdomain (e.g., cdf.toronto), and host (e.g., eddie).



- Key for primary domain = its name.
- Key for secondary domain = its name + name of primary domain.
- Key for host = its name + key of secondary domain = its name + name of secondary domain + name of primary domain.


All “Connecting” Entity Sets Are Weak




- If *Movies* determine *Studios*, we can omit *Studios name* from the key, and remove the double diamond from *Studios*.
- Better: *Studios* is attribute of *Contracts*.

Aggregation

Consider a relationship, *filmed_by*, between two entity sets, *Movies* and *Crews*




How can we add *Expenses* to the mix?

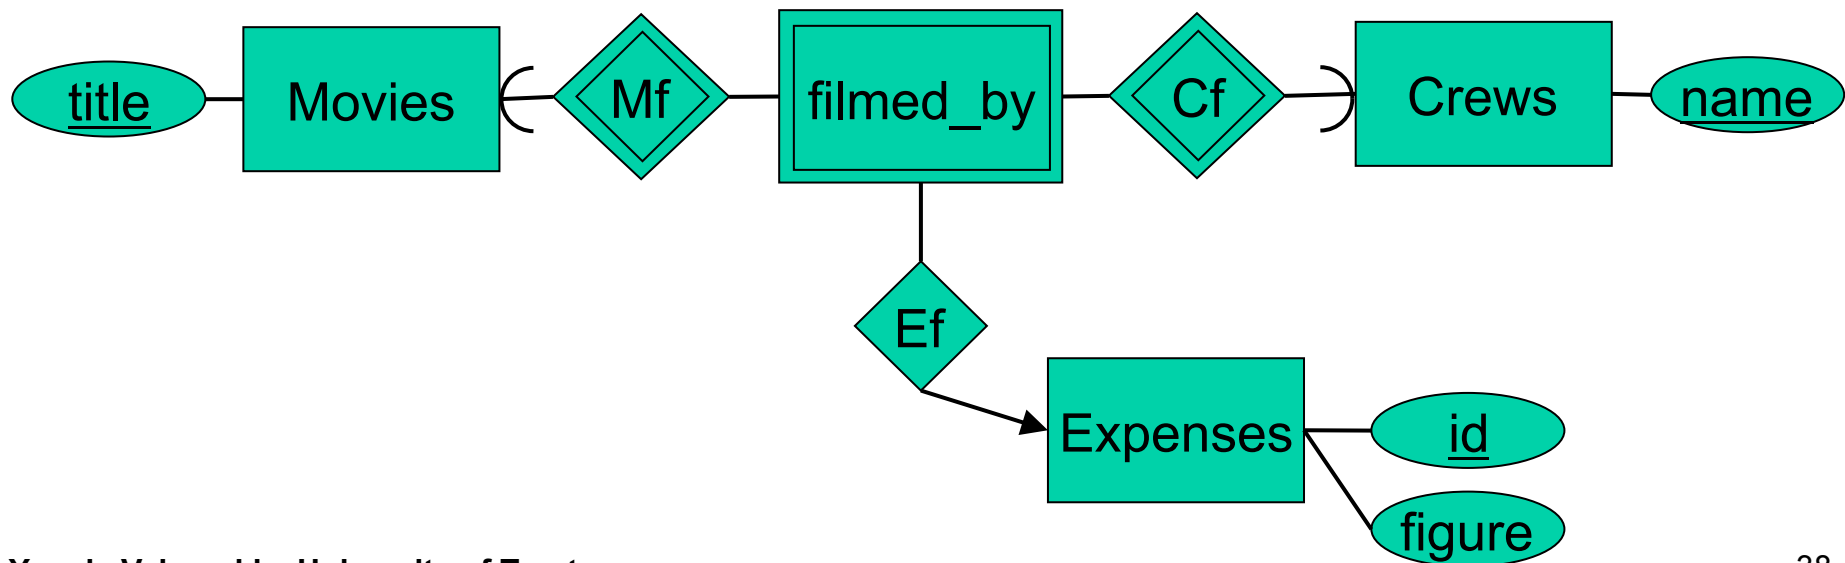


Simulating Aggregation

Solution: convert *filmed_by* into a weak entity set.



And then add *Expenses* to it...



Design Principles


Setting: client has (possibly vague) idea of what he/she wants. You must design a database that represents these thoughts and only these thoughts.

Important note:

Avoid redundancy, i.e., saying the same thing more than once. It wastes space and encourages inconsistency.


Example

Good:




Example

This is a bad design...



Yeap! That's also a bad design...



Use Schema to Enforce Constraints

The design *schema* should enforce as many constraints as possible.


- ◆ Don't rely on future data to follow assumptions.

Example:

- If a user must associate only one director with a movie, don't allow sets of directors and count on users to enter only one director per movie.

Entity Sets Vs. Attributes

- You may be unsure which concepts are worthy of being entity sets, and which are handled more simply as attributes.
- Especially tricky for the class design project, since there is a temptation to create needless entity sets to make project “larger.”



Which one do
you like???




Intuitive Rule for E.S. Vs. Attribute

Make an entity set only if it either:

1. Is more than a name of something; *i.e.*, it has non-key attributes or it has relationships with a number of different entity sets, or
2. Is the “many” in a many-one relationship.

Example

The following design illustrates both points:



- *Studios* deserves to be an E.S. because we record *addr*, a non-key attribute.
- *Movies* deserves to be an E.S. because it is at the “many” end.

Don't Overuse Weak E.S.

- There is a tendency to feel that no E.S. has its entities uniquely determined without following some relationships.
- However, in practice, we almost always create unique ID's to compensate: social-security numbers, OHIP's, etc.
- The only times weak E.S.'s seem necessary are when:
 - a) We can't easily create such ID's; e.g., no one is going to accept a “species ID” as part of the standard nomenclature (species is a weak E.S. supported by membership in a genus).
 - b) There is no global authority to create them, e.g., crews and studios.

The Relational Model

Relational Model

- Table = relation.
- Column headers = attributes.
- Row = tuple

Relation stars-in	
Movies	Stars
Basic Instinct	Sharon Stone
Total Recall	Arnold Schwarzenegger
Total Recall	Sharon Stone
...	...

- Relation schema = name(attributes) + other structure info., e.g., keys, other constraints.
Example: *stars-in(movie-name, star-name)*
 - ◆ Order of attributes is arbitrary, but in practice we need to assume the order given in the relation schema.
- Relation instance is current set of rows for a relation schema.
- Database schema is a collection of relation schemas.

Why Relations?


- Very simple model.
- *Often* a good match for the way we think about our data.
- Abstract model that underlies SQL, the most important language in DBMS's today.
 - ◆ But SQL uses “bags” while the abstract relational model is set-oriented.

Relational Design

We start with an E/R design. Then, the simplest approach (but not always best) is :

- convert each E.S. to a relation and
- convert each relationship to a relation.

Example:



Movies(title, year, length, film-type)

Note: E.S. attributes become relational attributes

Keys in Relations

An attribute (or set of attributes) K is a *key* for a relation R if we expect that in no instance of R will two different tuples agree on all the attributes of K . We indicate a key by underlining the key attributes. (pretty much what we did in E/R)

Example:

If *title* is a key for *Movies*, then

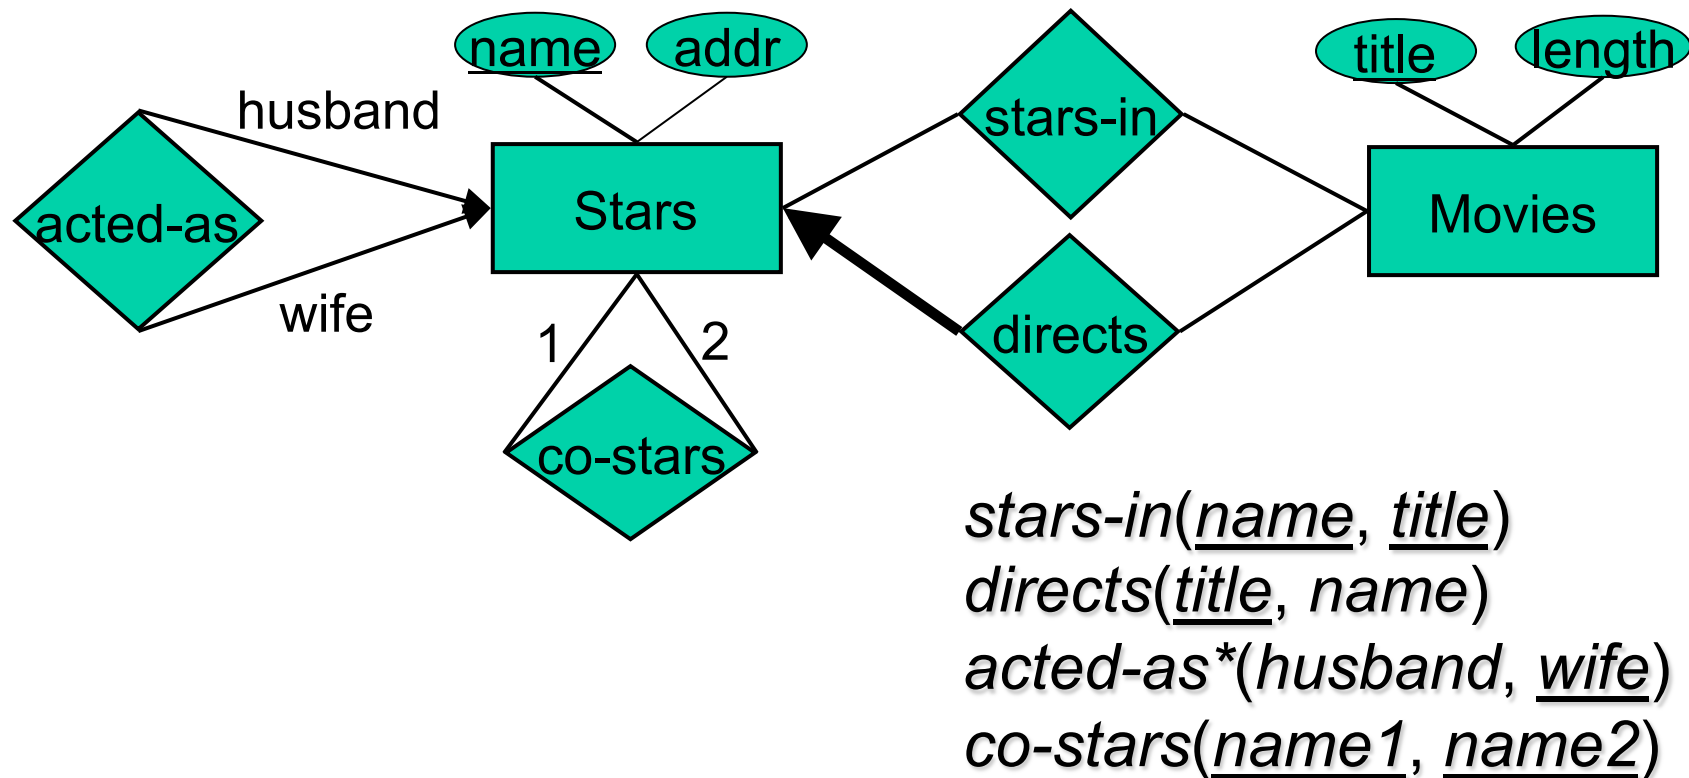
Movies(*title*, *year*, *length*, *film-type*)

From E/R Relationships to Relations

The constructed relation has the following attributes:

- The *key* attributes of each E.S. that participates in the relationship.
- Any attributes that belong to the relationship itself.
- Renaming attributes is necessary if an E.S. has multiple roles in the relationship

From E/R Relationships to Relations (cont.)



(*) For *acted-as*, we can choose either *husband* or *wife* as key

Combining Relations

Common case: Combine relation for an E.S. E with the relation for some many-one relationship from E to another E.S. F

Example: Combine *Movies*(title, length) with *directs*(title, name) to get *Movies-new*(title, length, director-name).

However danger in pushing this idea too far: redundancy


Example: Combine *Movies*(title, length) with *stars-in*(name, title) to get *Movies-stars*(title, name, length,)

<u>title</u>	<u>name</u>	<u>length</u>
Star Wars II	Ewan McGregor	142
Star Wars II	Natalie Portman	142

Weak Entities & Relations

- Relation for a weak E.S. must include its full key (*i.e.*, attributes of related entity sets) as well as its own attributes.
- A supporting (double-diamond) relationship yields a relation that is actually redundant and should be deleted from the database schema.

Weak Entity Sets & Relationships to Relations (cond.)



Hosts(name)
Logins(name, hname)
At(lname, hname, *hname2*)

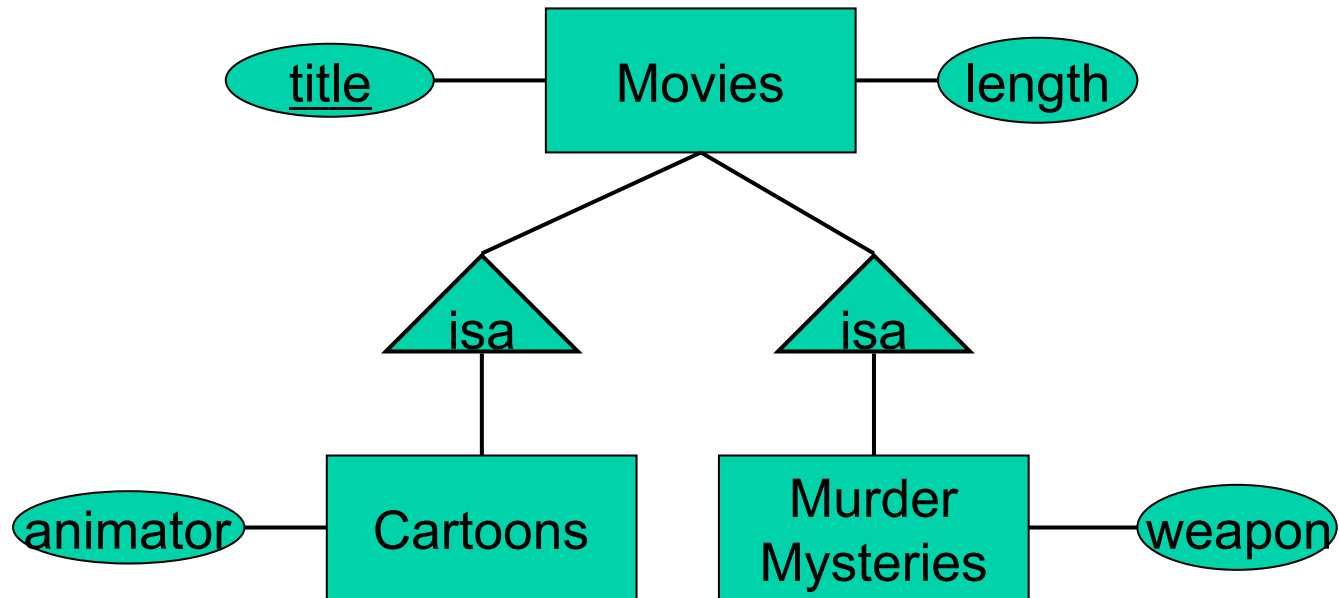
- In *At*, *hname* and *hname2* must be the same host, so delete one of them.
- Then, *Logins* and *At* become the same relation; delete one of them.
- In this case, *Hosts*' schema is a subset of *Logins*' schema. Delete *Hosts*?

Subclasses to Relations

Three approaches:

1. Object-oriented: each entity is in one class. Create a relation for each class, with all the attributes for that class.
 - ◆ Don't forget inherited attributes.
2. E/R style: an entity is in a network of classes related by *isa*. Create one relation for each E.S.
 - ◆ An entity is represented in the relation for each subclass to which it belongs.
 - ◆ Relation has only the attributes attached to that E.S. + key.
3. Use nulls. Create one relation for the root class or root E.S., with all attributes found anywhere in its network of subclasses.
 - ◆ Put NULL in attributes not relevant to a given entity.

Example



Subclasses to Relations: An Example

OO Style:

<u>Title</u>	<u>Length</u>	<u>Weapon</u>	<u>Title</u>	<u>Length</u>	<u>Animator</u>	<u>Title</u>	<u>Length</u>
Episode I	133	Lightsaber	Monsters	92	Yannis	MXXXXXX	92
...

E/R Style:

<u>Title</u>	<u>Length</u>	<u>Title</u>	<u>Weapon</u>	<u>Title</u>	<u>Animator</u>
Episode I	133	Episode I	Lightsaber	Monsters	Yannis
Monsters	92
...	...				

Using Nulls:

<u>Title</u>	<u>Length</u>	<u>Weapon</u>	<u>Animator</u>
Episode I	133	Lightsaber	Null
Monsters	92	Null	Yannis