
SQL: Databases Modification

Yannis Velegrakis

Database Modifications

Modification = insert + delete + update

e.g. to insert a tuple we use the statement

INSERT INTO <relation> **VALUES** (<list of values>)

- In the above statement the values as listed in the same order with which the attributes were declared
- If we want to ignore this order then we list the attributes as arguments of the relation

Example:

Consider *StarsIn*(movieTitle, movieYear, starName)

Insert the fact that Sydney Green stars in The Maltese Falcon

INSERT INTO *StarsIn*(*movieTitle*, *starName*, *movieYear*)
VALUES('The Maltese Falcon', 'Sydney Green', 1942);

Insertion of the Result of a Query

Syntax: **INSERT INTO** <relation> (<subquery>).

Example:

Consider

Studio(*name*, *address*, *presC#*)

Movie(*title*, *year*, *length*, *inColor*, *studio-Name*, *prodC#*)

Add to relation *Studio* all the studios that are mentioned in relation *Movie*

```
INSERT INTO Studio(name)  
  SELECT DISTINCT studio-Name  
  FROM Movie  
  WHERE studio-Name NOT IN (SELECT name  
                                FROM Studio)
```

Deletion

Syntax: **DELETE FROM** <relation> **WHERE** <condition>

Semantic: Deletes all tuples satisfying the condition from
the named relation

Example:

Consider *StarsIn*(movieTitle, movieYear, starName)

Sydney Green was not a star in The Maltese Falcon...

```
DELETE FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon' AND
      movieYear = 1942 AND starName = 'Sydney Green'
```

As another example, to make the *StarsIn* relation empty execute

```
DELETE FROM StarsIn
```

Updates

Syntax: **UPDATE** <relation>

SET <new-value assignments>

WHERE <condition>

Example:

Consider

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

Prefix the name of every movie executive with the title 'Pres. '
if the executive is the president of a studio.

UPDATE *MovieExec*

SET *name* = 'Pres. ' || *name*

WHERE *cert#* IN (**SELECT** *presC#* **FROM** *Studio*)

Defining a Database Schema

Syntax: **CREATE TABLE** <name> (<list of elements>)

- Principal elements are attributes and their types, but key declarations and constraints also appear.
- Similar **CREATE X** commands for other schema elements **X**: **VIEWS**, **INDEX**, **ASSERTION**, **TRIGGER**.
- “**DROP X name**” deletes the created element of kind **X** with that name.

Example: *Studio(name, address, presC#)*

```
CREATE TABLE Studio (  
    name CHAR(20),  
    address VARCHAR(20),  
    presC# REAL)
```

```
DROP TABLE Studio
```

Types

1. **INT** or **INTEGER**.
2. **BOOLEAN**
3. **REAL** or **FLOAT**.
4. **CHAR**(n) = fixed length character string, padded with “pad characters.”
5. **VARCHAR**(n) = variable-length strings up to n characters.
6. **DATE & TIME**

Declaring Keys

Use **PRIMARY KEY** or **UNIQUE**.

- But only one primary key, many **UNIQUE**s allowed.
- SQL permits implementations to create an *index* (data structure to speed access given a key value) in response to **PRIMARY KEY** only.
- SQL does not allow nulls in primary key, but allows them in “unique” columns (which may have two or more nulls, but not repeated non-null values).

Declaring Keys

Two places to declare:

1. After an attribute's type, if the attribute is a key by itself.
2. As a separate element.
 - ◆ Essential if key is >1 attribute.

Example

```
CREATE TABLE Studio (  
    name CHAR(20),  
    address VARCHAR(20),  
    presC# REAL,  
    PRIMARY KEY(name))
```

Example

Find the differences...

```
CREATE TABLE Studio-version1 (  
    name CHAR(20),  
    address VARCHAR(20) UNIQUE,  
    presC# REAL UNIQUE,  
    PRIMARY KEY(name))
```

```
CREATE TABLE Studio-version2 (  
    name CHAR(20),  
    address VARCHAR(20),  
    presC# REAL,  
    UNIQUE (address, presC#),  
    PRIMARY KEY(name))
```

Other Properties You Can Give to Attributes

1. **NOT NULL** = every tuple must have a real value for this attribute.
2. **DEFAULT** value = a value to use whenever no other value of this attribute is known.

Example:

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St',  
    presC# REAL NOT NULL)
```

Example

Consider executing the following statement

```
INSERT INTO Studio(name, presC# )  
VALUES('Fox', 90943)
```

The result is the following tuple:

<i>name</i>	<i>addr</i>	<i>presC#</i>
Sally	123 Sesame St.	90943

Primary key is by default not NULL.

This insertion is legal. It is OK to list a subset of the attributes and values for only this subset. But if we had forgot to specify a value for *presC#* then the insertion could not be made.

Changing Columns

Add an attribute to relation R with

ALTER TABLE R ADD <column declaration>

Example:

```
CREATE TABLE MovieStar (  
    name CHAR(30),  
    address VARCHAR(255),  
    gender CHAR(1),  
    birthdate DATE,  
    PRIMARY KEY(name)
```

```
ALTER TABLE MovieStar ADD phone CHAR(16) DEFAULT 'unlisted'
```

Columns may also be dropped...

Example"

```
ALTER TABLE MovieStar DROP birthdate
```

Constraints

Commercial relational systems allow much more “fine-tuning” of constraints than do the modeling languages we learned earlier.

In essence: SQL programming is used to describe constraints.

Outline:

1. Primary key declarations (already covered).
2. Foreign-keys = referential integrity constraints.
3. Attribute- and tuple-based checks = constraints within relations.
4. SQL Assertions = global constraints

Foreign Keys

In relation R a clause that “attribute A **REFERENCES** $S(B)$ ” says that whatever values appear in the A column of R must also appear in the B column of relation S .

NOTE: B must be declared as the primary key for S .

Example:

Consider

MovieExec(*name*, *address*, *cert#*, *netWorth*)

Studio(*name*, *address*, *presC#*)

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St',  
    presC# REAL REFERENCES MovieExec(cert#))
```


Foreign Keys (cont.)

Alternative:

Add another element declaring the foreign key, as:

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St',  
    presC# REAL.  
    FOREIGN KEY (presC#) REFERENCES MovieExec(cert#))
```

Extra element is essential if the foreign key has
more than one attribute...

Foreign Key Constraint Violation

Two ways to violate the constraint:

1. Insert or update a *Studio* tuple so it refers to a non-existent movie executive
 - ◆ Always rejected.
2. Delete or update a *MovieExec* tuple that has a *cert#* value some *Studio* tuples refer to.
 - a) Default: reject.
 - b) *Cascade*: Ripple changes to referring *Studio* tuples.
 - c) *Set Null*: Change referring tuples to have **NULL** in referring components.

Selecting a Policy

Add **ON [DELETE, UPDATE] [CASCADE, SET NULL]** to the declaration of foreign key.

Example:

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St',  
    presC# REAL.  
    FOREIGN KEY (presC#) REFERENCES MovieExec(cert#)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE)
```

Choosing a “correct” policy is a design decision...

Attribute-Based Checks

Follow an attribute by a condition that must hold for that attribute in each tuple of its relation.

- Syntax: **CHECK** (<condition>).
 - ◆ Condition may involve the checked attribute.
 - ◆ Other attributes and relations may be involved, but *only* in subqueries.
- Condition is checked only when the associated attribute changes (*i.e.*, an insert or update occurs).

Example

```
CREATE TABLE Studio (  
  name CHAR(30) PRIMARY KEY,  
  addr CHAR(50) DEFAULT '123 Sesame St',  
  presC# REAL CHECK (presC# >= 100000),  
  FOREIGN KEY (presC#) REFERENCES MovieExec(cert#)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE)
```

- Check on *presC#* is like a foreign-key constraint, except:
 - ◆ The check occurs only when we add a tuple or change the *presC#* in an existing tuple, not when we delete a tuple from *Studio*

Tuple-Based Checks

Separate element of table declaration.

- Syntax: like attribute-based check.
- But condition can refer to any attribute of the relation.
 - ◆ Or to other relations/attributes in subqueries.
- Checked whenever a tuple is inserted or updated.

Example

If a star's gender is male, then his name must not begin with 'Ms.'

```
CREATE TABLE MovieStar (  
    name CHAR(30) PRIMARY KEY,  
    address CHAR(255),  
    gender CHAR(1),  
    birthdate DATE,  
    CHECK (gender = 'F' OR name NOT LIKE 'Ms. %')
```

SQL Assertions

- Database-schema constraints
- Checked whenever a mentioned relation changes.
- Syntax:

```
CREATE ASSERTION < name>  
CHECK(<condition>)
```


Example

Consider

MovieExec(*name*, *address*, *cert#*, *netWorth*)

Studio(*name*, *address*, *presC#*)

What if we wish to require that no one can become the president of a studio unless their net worth is at least \$10,000,000

```
CREATE ASSERTION RichPres CHECK  
  (NOT EXISTS(SELECT *  
                FROM Studio, MovieExec  
                WHERE presC# = cert# AND  
                    netWorth < 10000000))
```

Checked whenever *Studio* and *MovieExec* change

Triggers

Often called event-condition-action rules.

- *Event* = a class of changes in the DB, e.g., “insertions into Beers.”
- *Condition* = a test as in a where-clause for whether or not the trigger applies.
- *Action* = one or more SQL statements.
- Differ from checks or SQL assertions in that:
 1. Triggers invoked by the event; the system doesn't have to figure out when a trigger could be violated.
 2. Condition not available in checks.

Example

Consider

MovieExec(name, address, cert#, netWorth)

Whenever we update the netWorth of a movie executive we want the new value to be bigger than the old value.

```
CREATE TRIGGER NetWorthTrigger
AFTER UPDATE OF netWorth ON MovieExec
REFERENCING
    OLD ROW AS OldTuple,
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN(OldTuple.netWorth > NewTuple.netWorth)
    UPDATE MovieExec
    SET netWorth = OldTuple.netWorth
    WHERE cert# = NewTuple.cert#
```